“Data Mining with Rattle and R”

DIT Analytics Club
07 March, 2013
“Peadar Kearney's”, Dame Street

Colman McMahon
colman@thedata.co

LinkedIn
www.linkedin.com/in/colmanmcmahon

Agenda

- Intro
- Data mining
- Rattle installation
- Rattle workflow
- Appendix

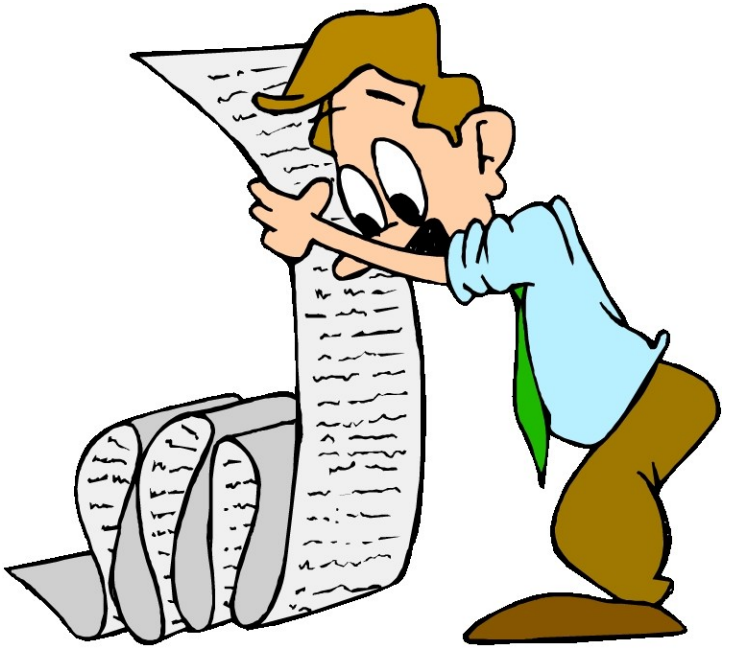

Agenda

- Intro
- Data mining
- Rattle installation
- Rattle workflow
- Appendix

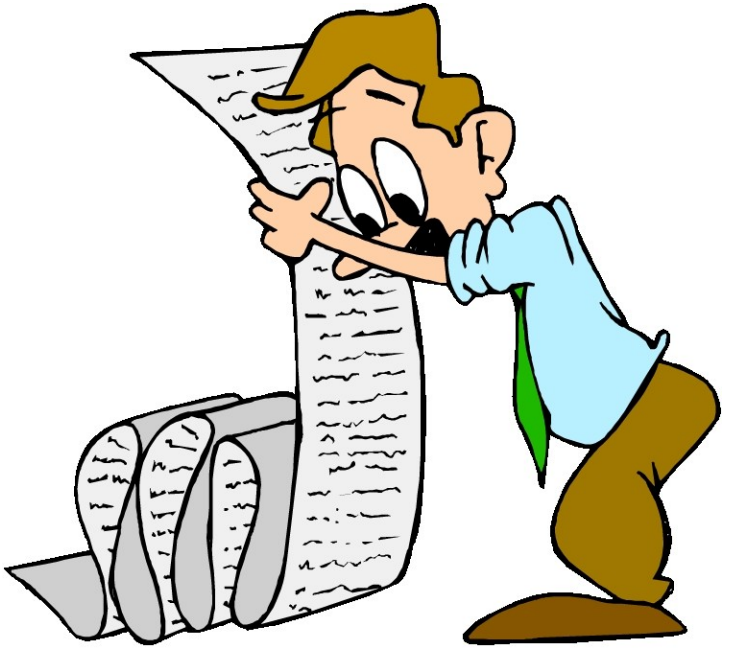

Bio

- Professional
  - » Background: Film, VFX, Digital Media
  - » Currently: PhD Fellow, UCD Dynamics Lab
- PhD:
  - » Policy network analysis
  - » Creative Industries
- Technologies:
  - » Social Network Analysis
  - » Agent Based Modelling
  - » Data/statistical analysis

www.dl.ucd.ie

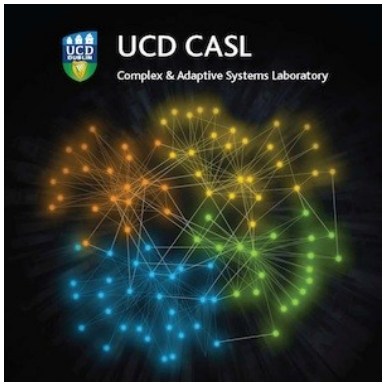

www.ucd.ie/casl/

- Data Mining with Rattle and R
  - » The Art of Excavating Data for Knowledge Discovery
- Graham Williams
- Springer, 2011
- ISBN 978-1-4419-9889-7

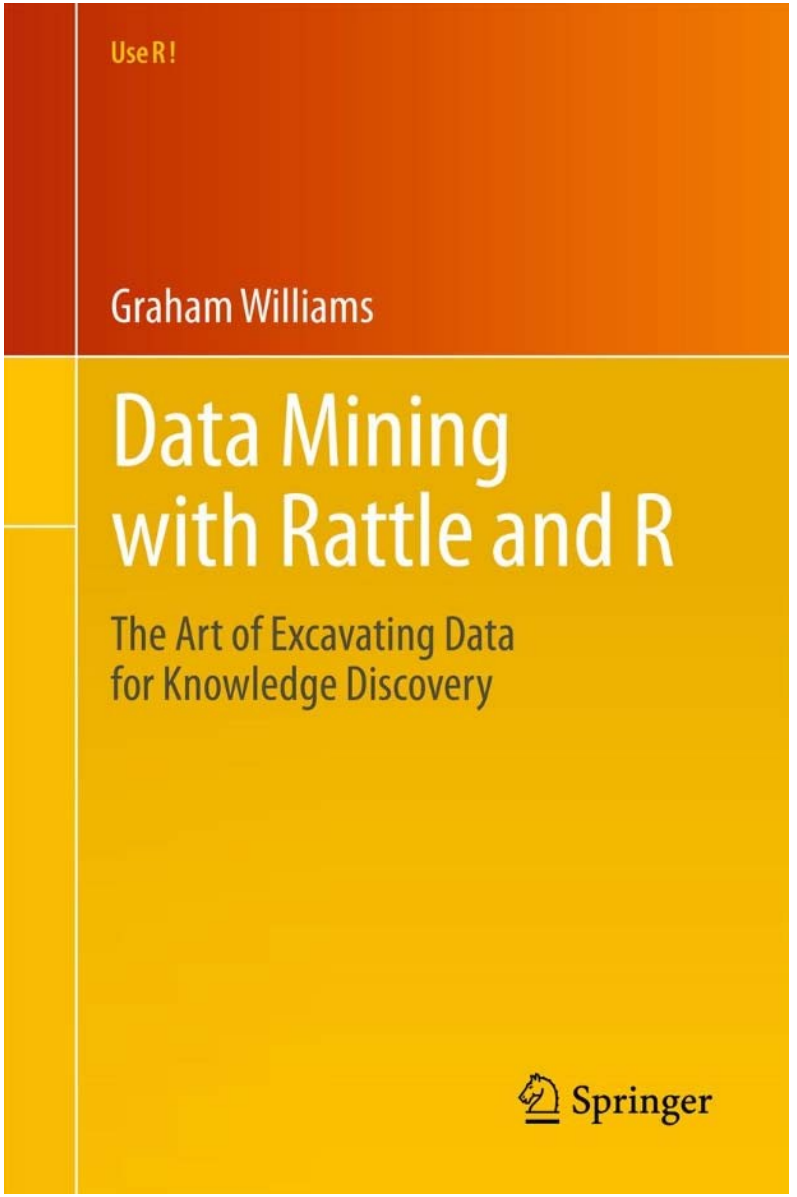

Rattle

togaware

Rattle - "the R Analytical Tool To Learn Easily"

- Presents statistical & visual summaries of data
- Transforms data into forms that can be readily modelled
- Builds models (unsupervised and supervised) from the data
- Graphically presents the performance of models
- Scores new datasets

Rattle

- Built on the statistical language R
  - » an understanding of R is not required in order to use it
- Simple to use, quick to deploy, and allows us to rapidly work through the data processing, modelling, and evaluation phases of a data mining project
- Can migrate from Rattle to R when we need to fine-tune and further develop our data mining projects

Rattle GUI

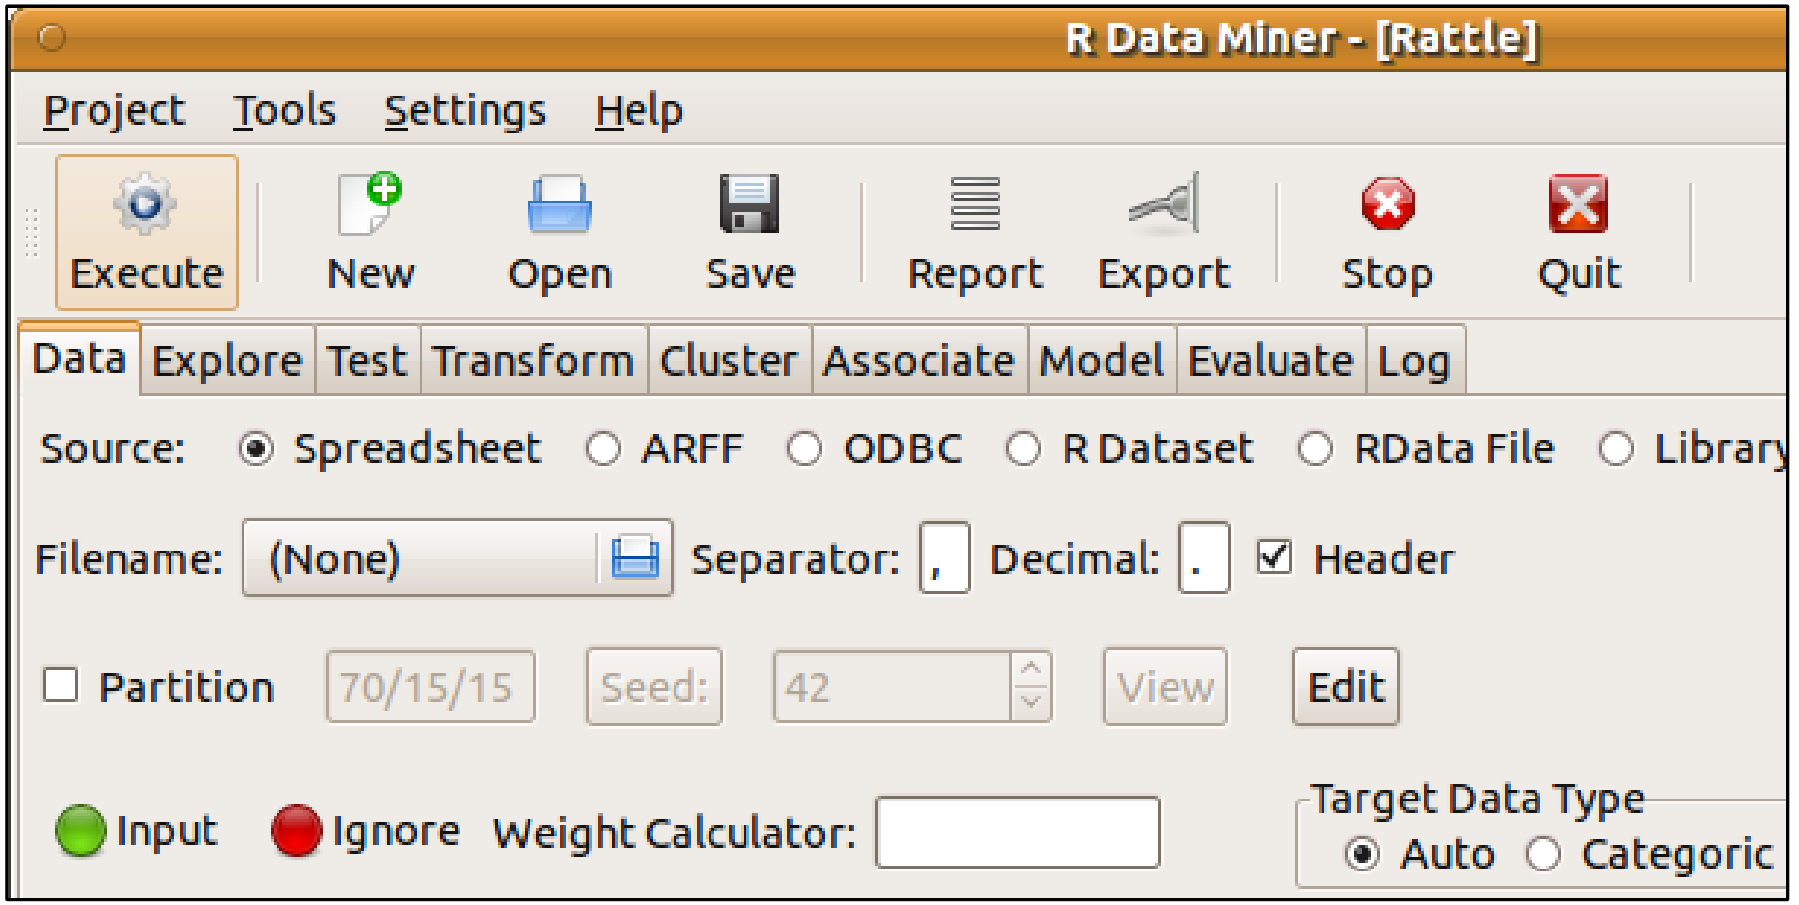
R Data Miner - [Rattle]
Project Tools Settings Help
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Source: Spreadsheet ARFF ODBC R Dataset RData File Library
Filename: (None) Separator: , Decimal: . Header
Partition 70/15/15 Seed: 42 View Edit
Input Ignore Weight Calculator:
Target Data Type
Auto Categoric

Agenda

- Intro
- Data mining
- Rattle stack
- Rattle workflow
- Appendix

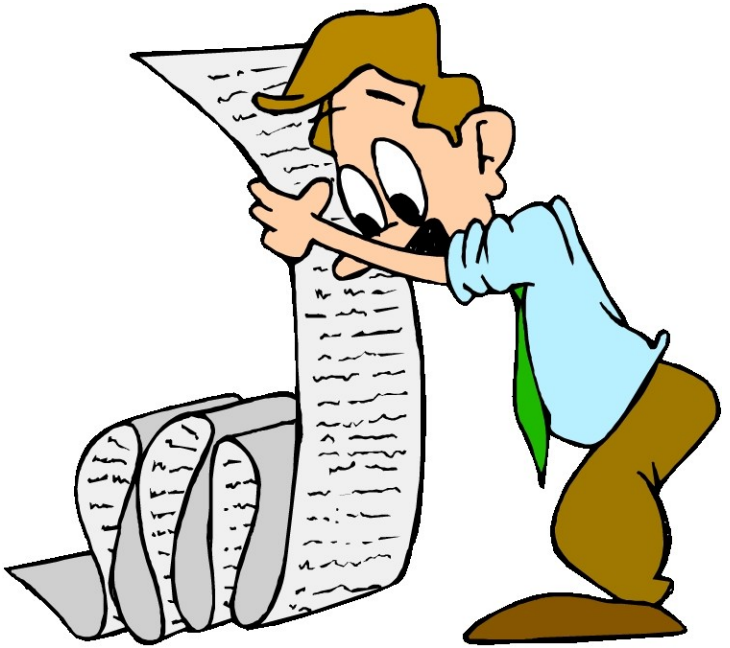

CRISP-DM

- Cross Industry Process for Data Mining (CRISP-DM, 1996)
  - » framework for delivering data mining projects.

1) **Problem Understanding**
2) **Data Understanding**
3) **Data Preparation**
4) **Modelling**
5) **Evaluation**
6) **Deployment**

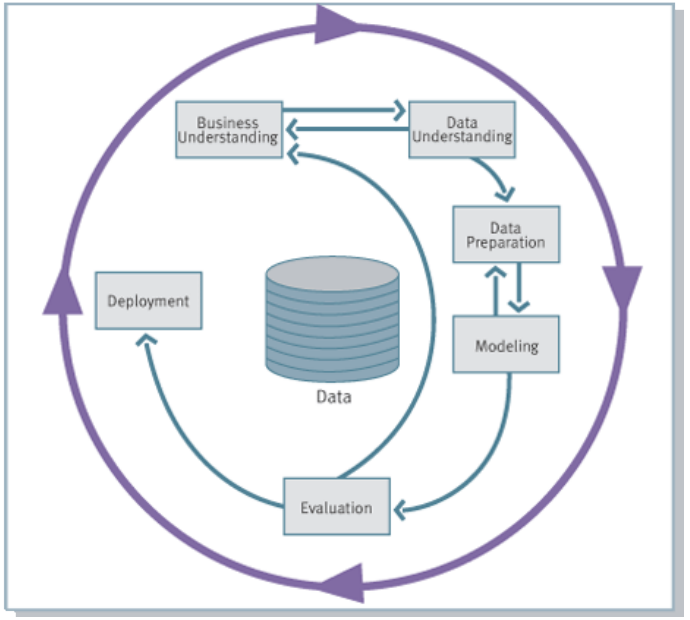

Rattle GUI

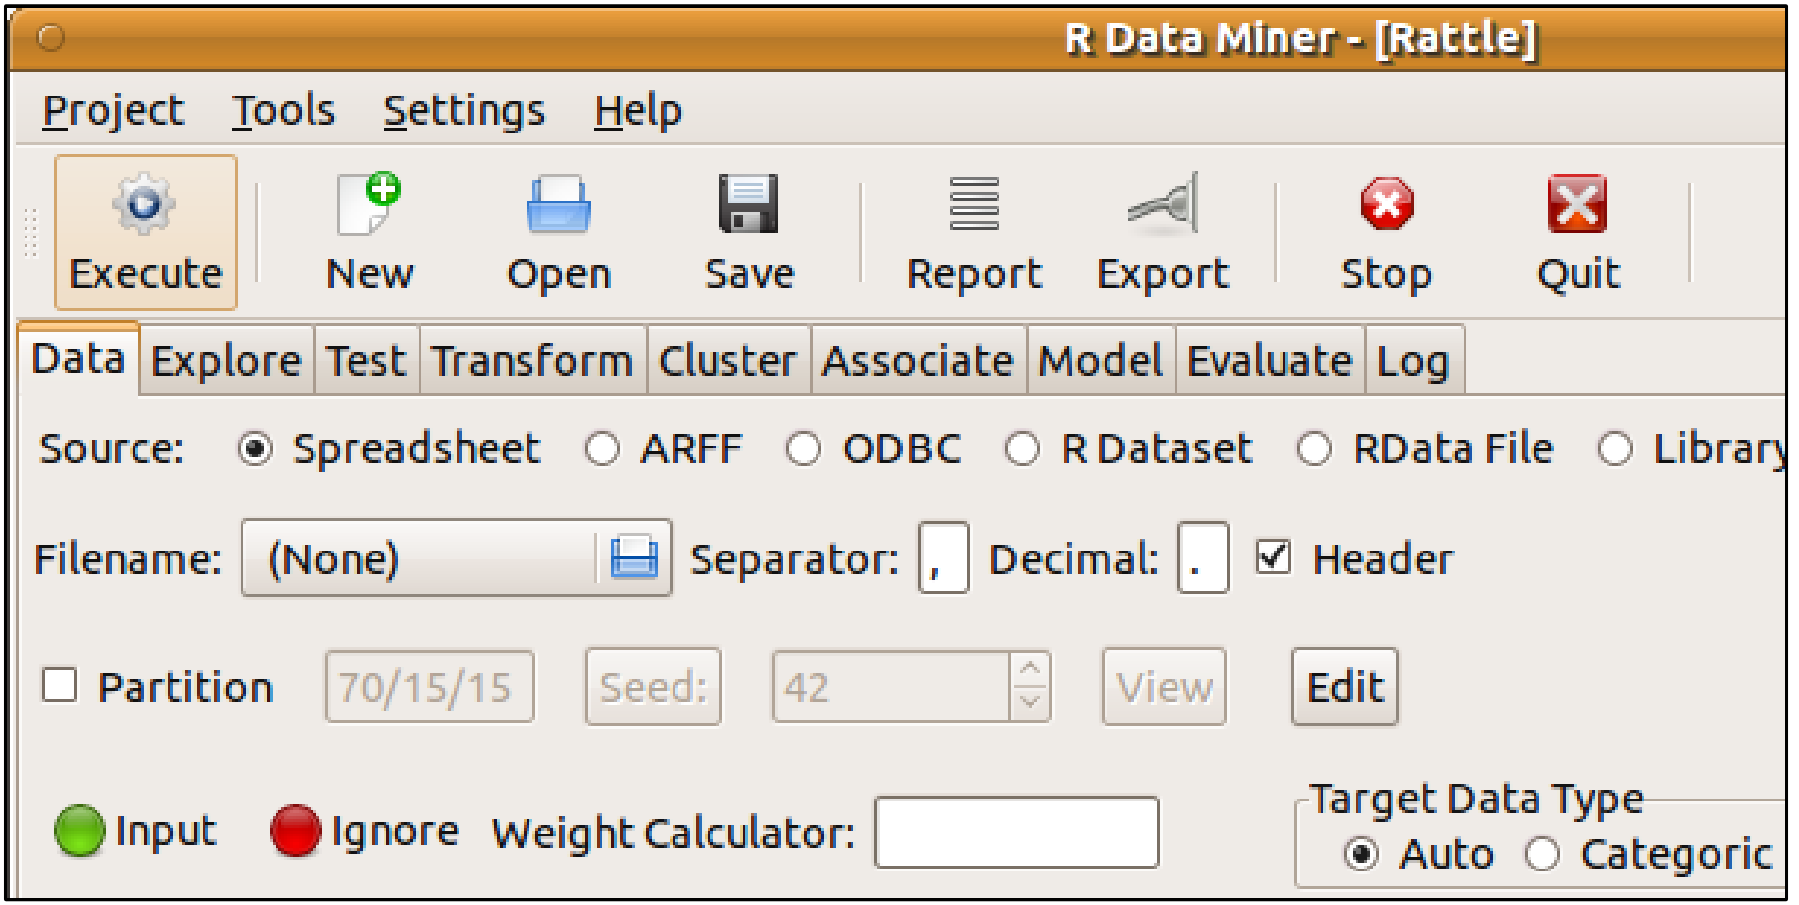
R Data Miner - [Rattle]
Project Tools Settings Help
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Source: Spreadsheet ARFF ODBC R Dataset RData File Library
Filename: (None) Separator: , Decimal: . Header
Partition 70/15/15 Seed: 42 View Edit
Input Ignore Weight Calculator:
Target Data Type
Auto Categoric

Agenda

- Intro
- Data mining
- Rattle installation
- Rattle workflow
- Appendix

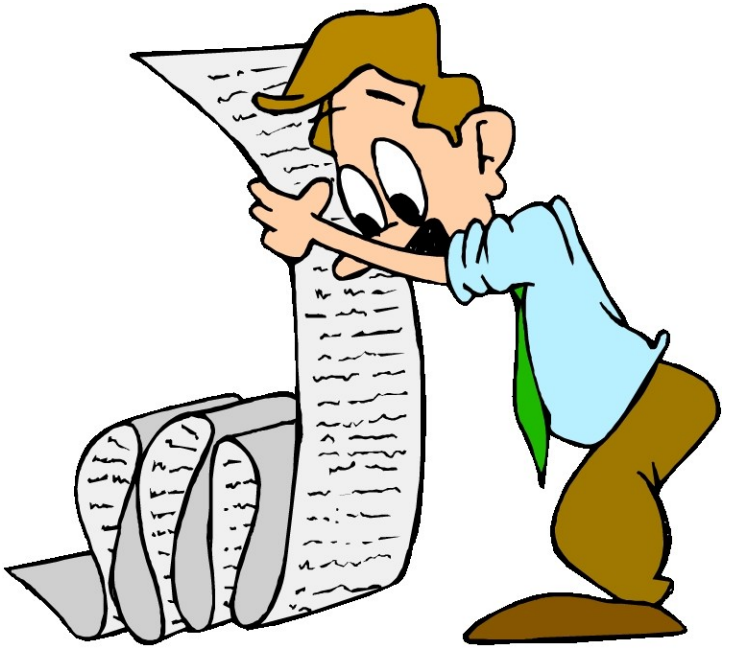

Rattle stack

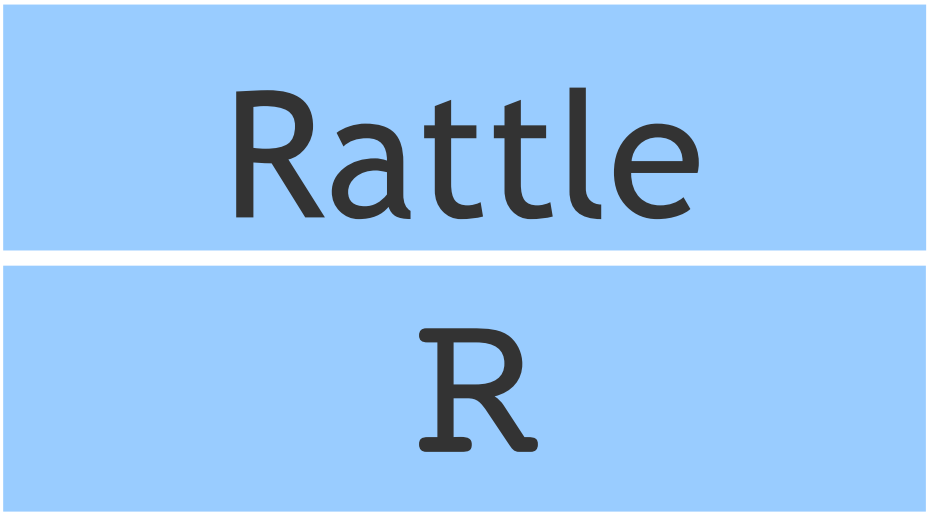

R (www.r-project.org)

- R - a sophisticated statistical software package
  - » easily installed, instructional, state-of-the-art, and it is free and open source
- The basic *modus operandi* - write scripts using the R language
- Steeper learning curve than using GUI based systems, but once over the hurdle, becomes relatively easy

R Project (www.r-project.org)

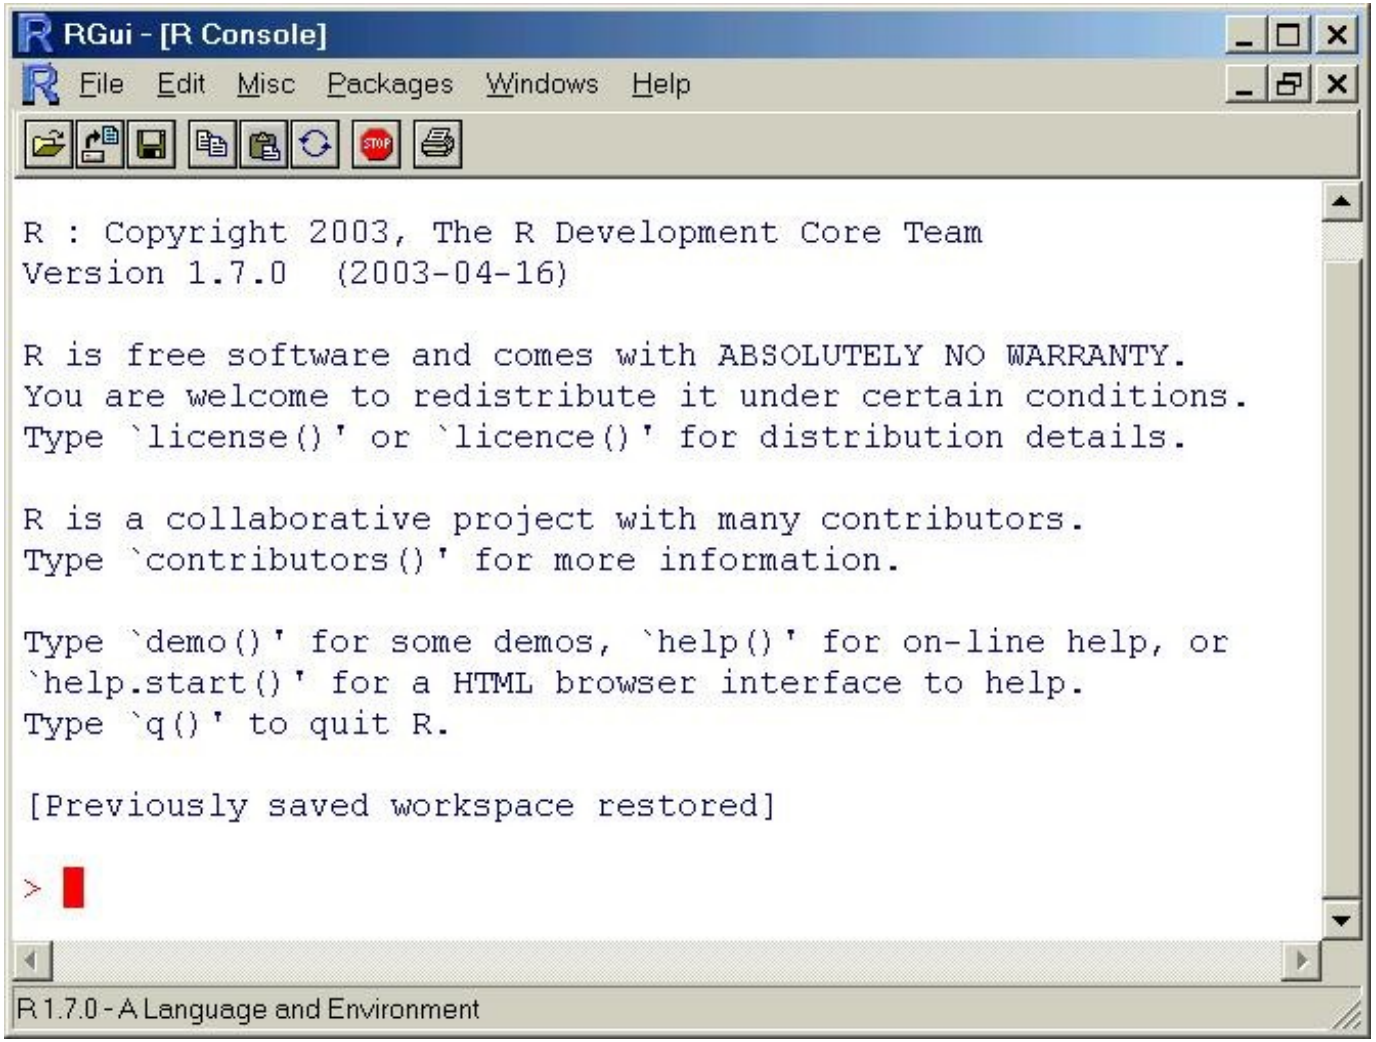

Rattle (rattle.togware.com)

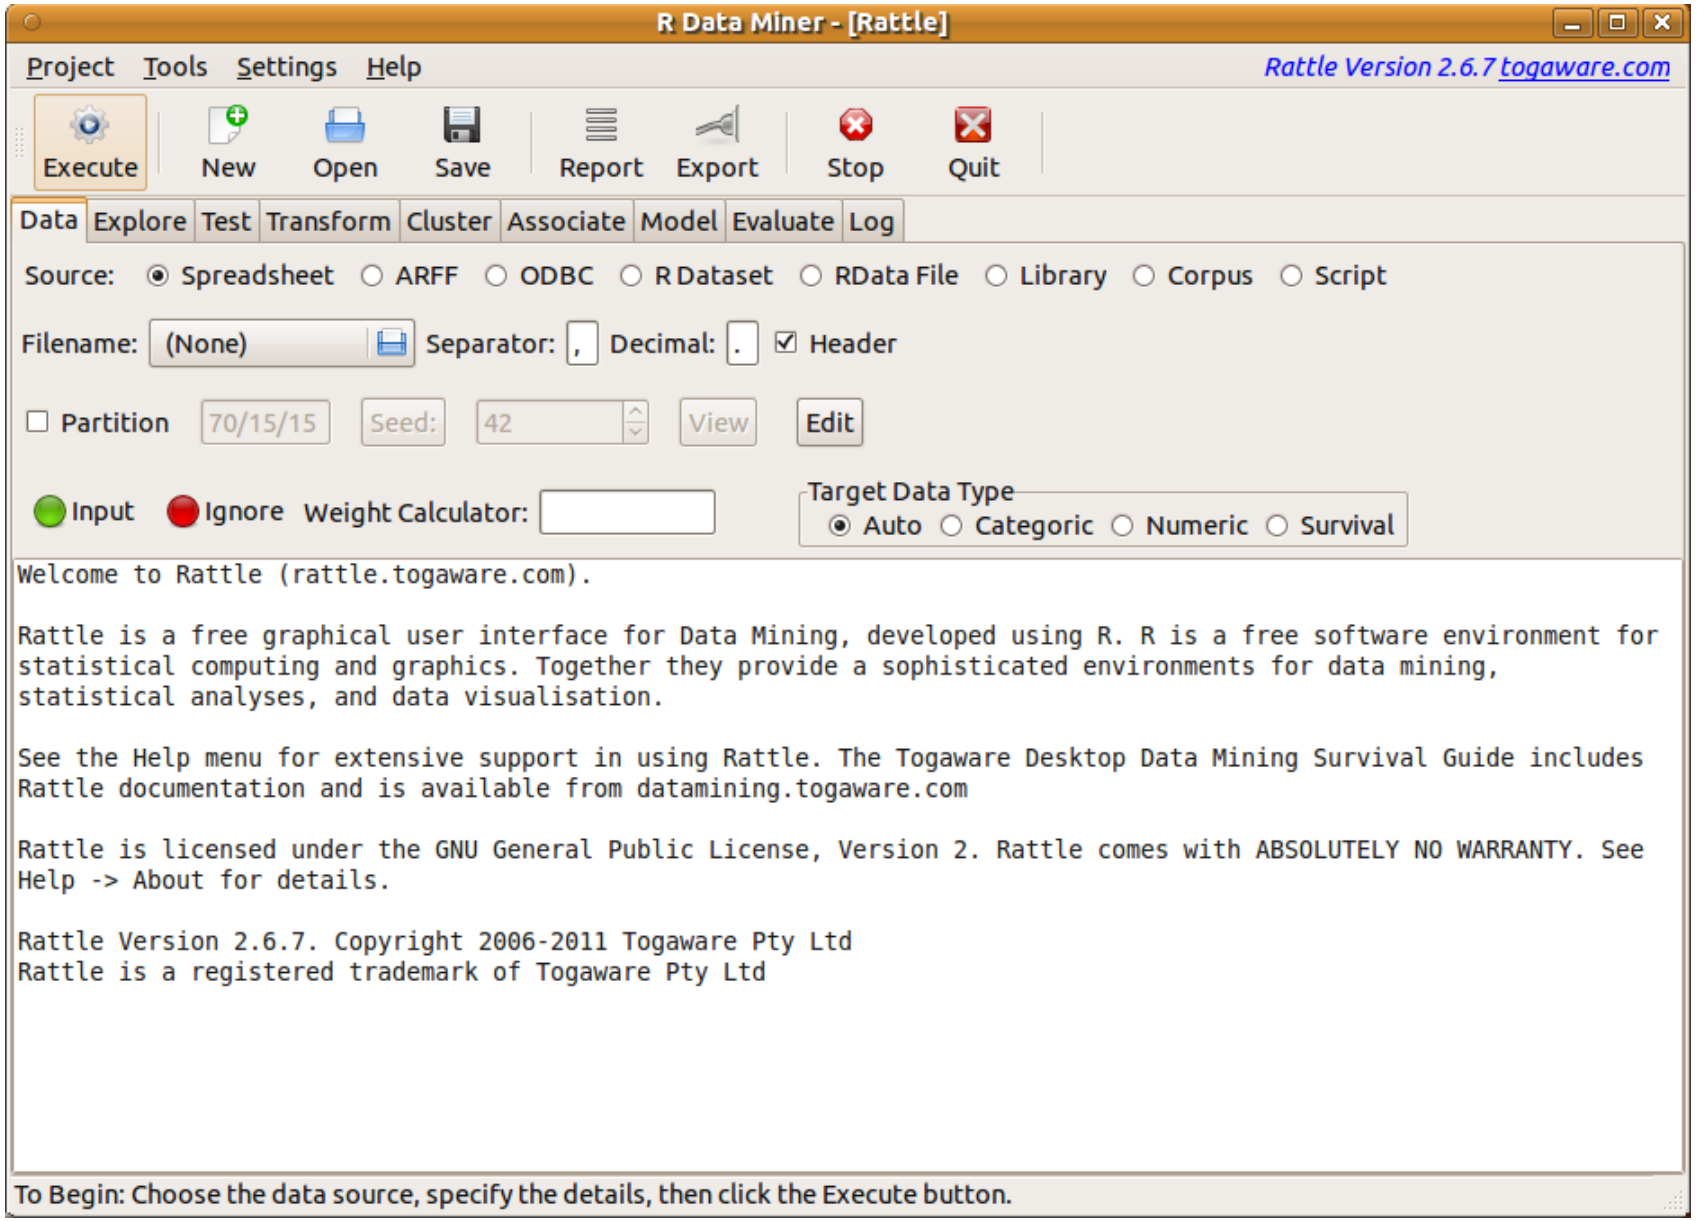

Installation

Install R

```
www.r-project.org
```

```
Start R
```

Install Rattle

```
> install.packages("rattle")
```

Load rattle into the R library

```
> library(rattle)
> rattle()
```

Agenda

- Intro
- Data mining
- Rattle installation
- Rattle workflow
- Appendix

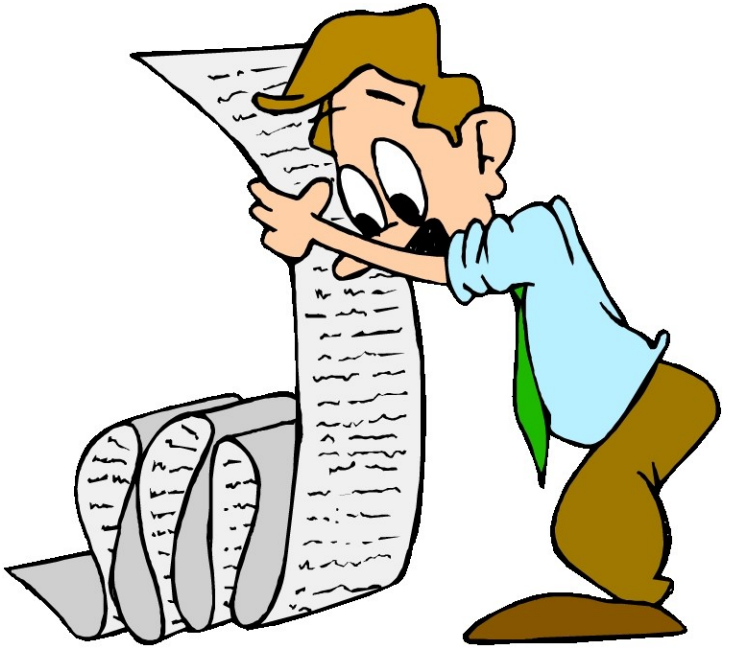

Rattle workflow

1) Load a dataset
2) Select variables and entities for exploring and mining
3) Explore the data to understand how it is distributed or spread
4) Transform the data to suit our data mining purposes
5) Build our models
6) Evaluate the models on other datasets
7) Export the models for deployment

Workflow - tabs

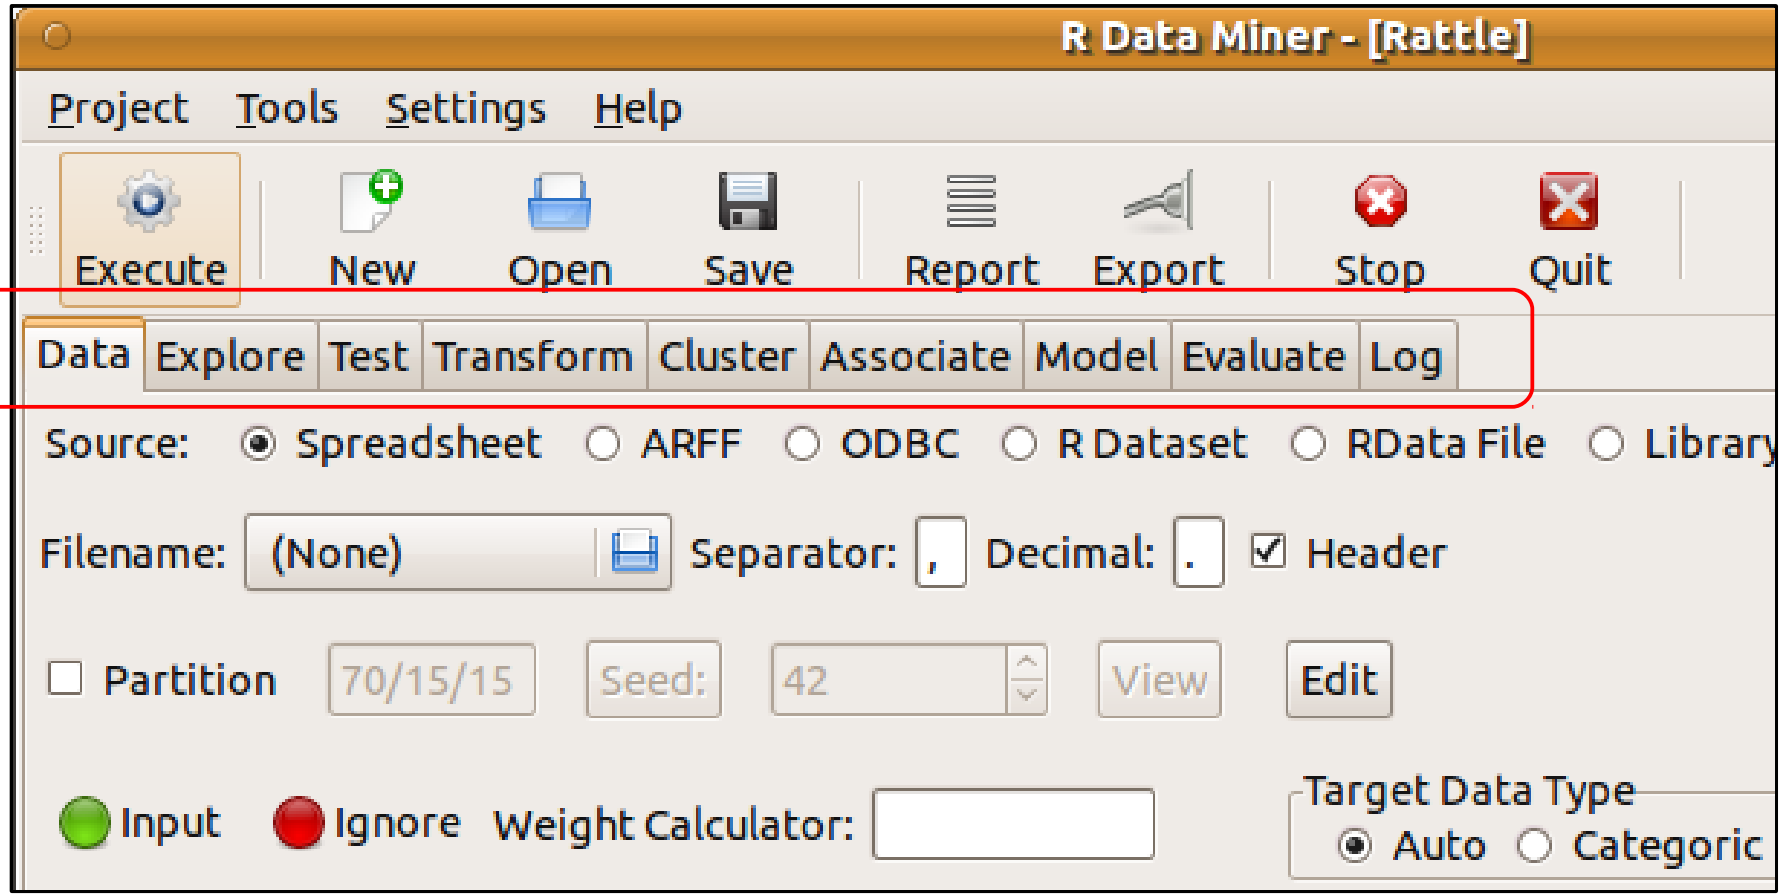

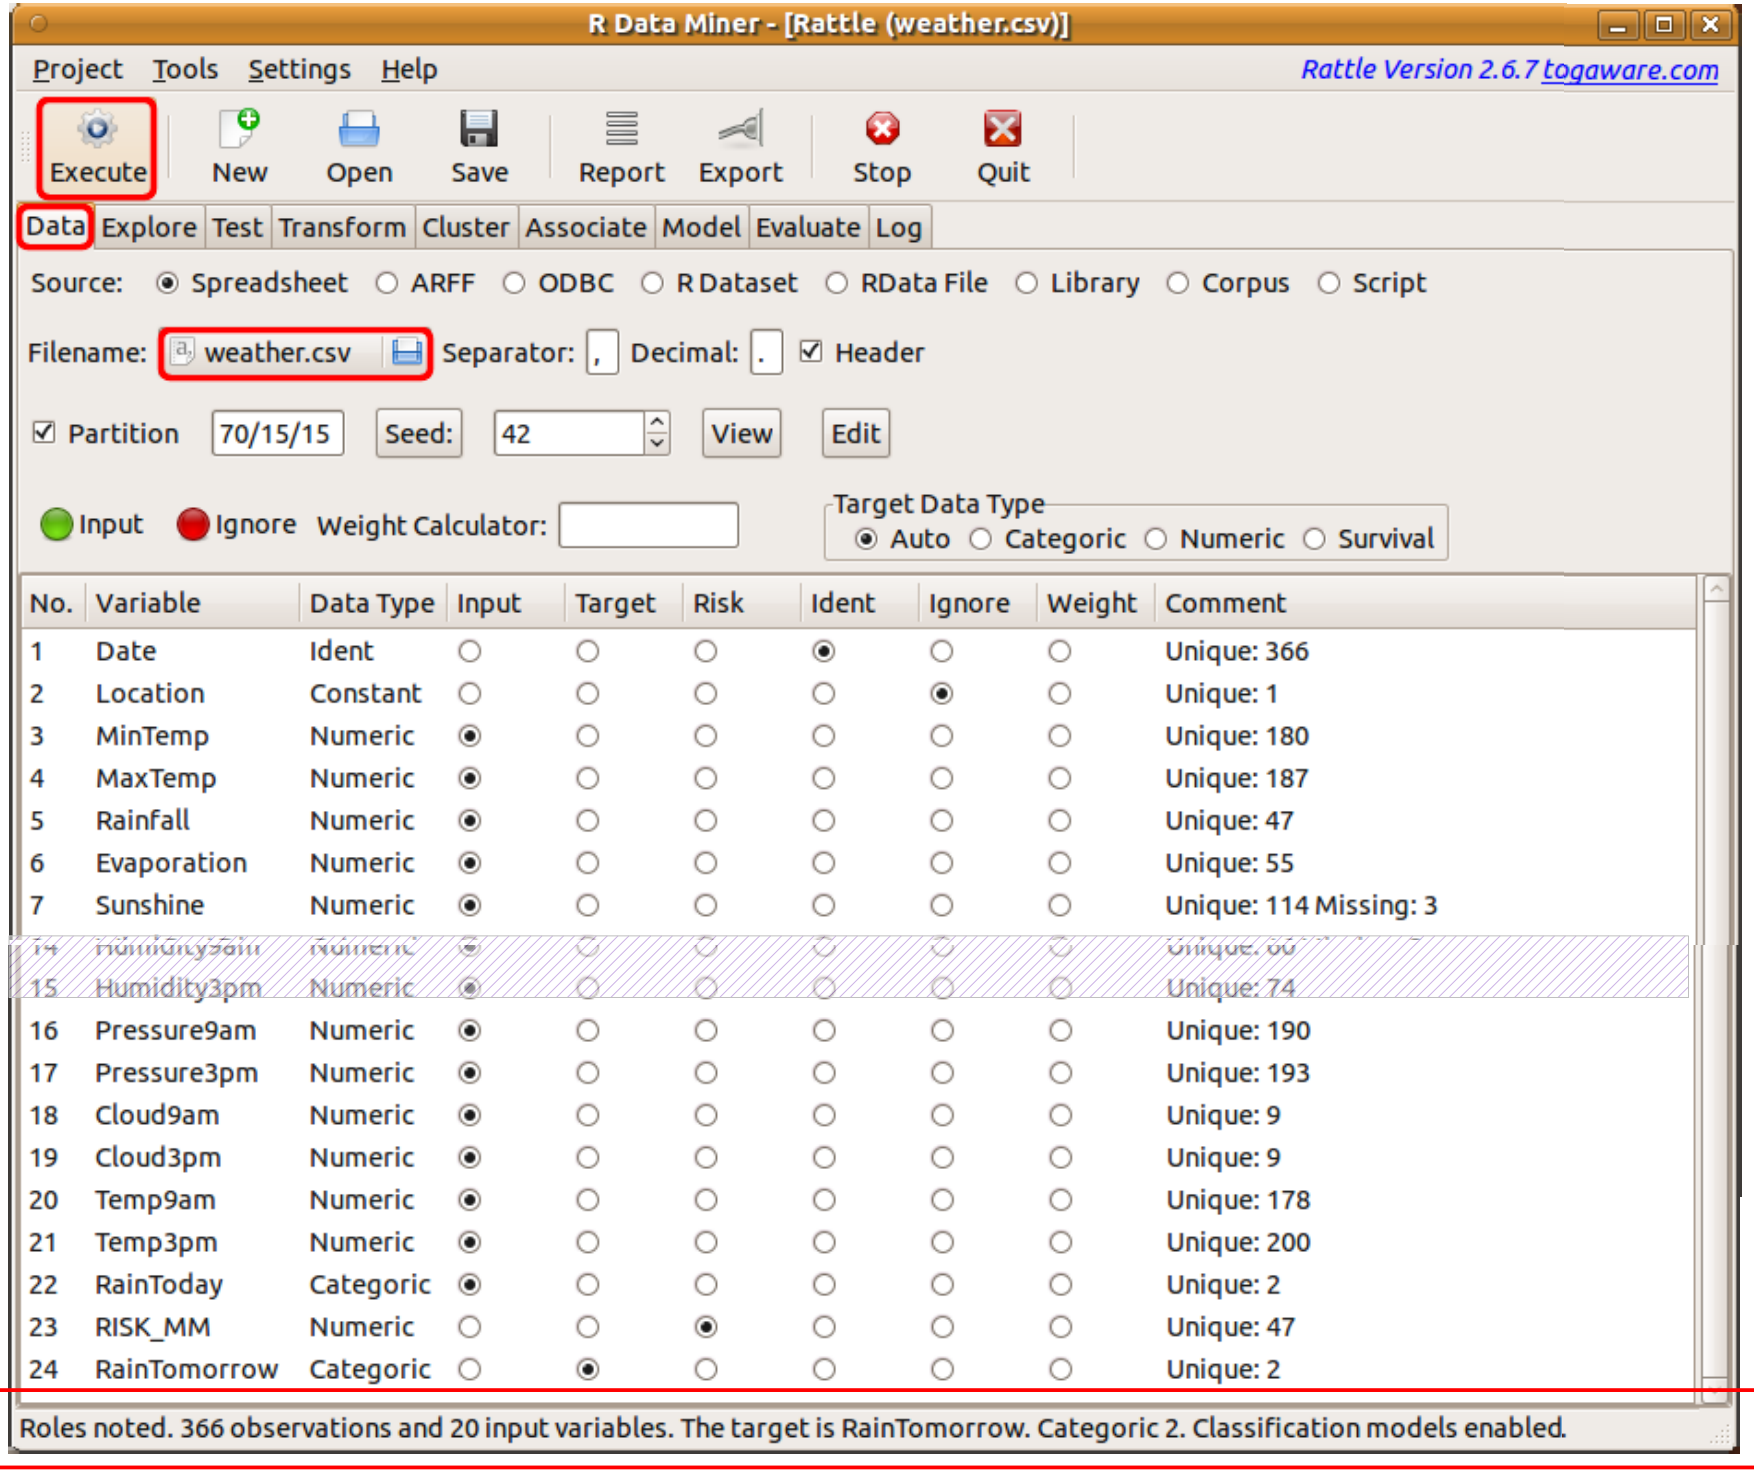
R Data Miner - [Rattle (weather.csv)]
Project Tools Settings Help
Rattle Version 2.6.7 togaware.com
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Source: Spreadsheet ARFF ODBC R Dataset RData File Library Corpus Script
Filename: weather.csv Separator: , Decimal: . Header
Partition 70/15/15 Seed: 42 View Edit
Input Ignore Weight Calculator:
Target Data Type
Auto Categoric Numeric Survival
No. Variable Data Type Input Target Risk Ident Ignore Weight Comment
1 Date Ident Unique: 366
2 Location Constant Unique: 1
3 MinTemp Numeric Unique: 180
4 MaxTemp Numeric Unique: 187
5 Rainfall Numeric Unique: 47
6 Evaporation Numeric Unique: 55
7 Sunshine Numeric Unique: 114 Missing: 3
15 Humidity3pm Numeric Unique: 74
16 Pressure9am Numeric Unique: 190
17 Pressure3pm Numeric Unique: 193
18 Cloud9am Numeric Unique: 9
19 Cloud3pm Numeric Unique: 9
20 Temp9am Numeric Unique: 178
21 Temp3pm Numeric Unique: 200
22 RainToday Categoric Unique: 2
23 RISK_MM Numeric Unique: 47
24 RainTomorrow Categoric Unique: 2
Roles noted. 366 observations and 20 input variables. The target is RainTomorrow. Categoric 2. Classification models enabled.

Explore


```
> summary(weather[7:9])

    Sunshine       WindGustDir   WindGustSpeed
 Min.   : 0.00   NW     : 73   Min.   :13.0
 1st Qu.: 5.95   NNW    : 44   1st Qu.:31.0
 Median : 8.60   E      : 37   Median :39.0
 Mean   : 7.91   WNW    : 35   Mean   :39.8
 3rd Qu.:10.50   ENE    : 30   3rd Qu.:46.0
 Max.   :13.60   (Other):144   Max.   :98.0
 NA's   : 3.00   NA's   :  3   NA's   : 2.0
```

- Provides a textual overview of the data

Explore - detailed


```
> library(fBasics)
> basicStats(weather$Sunshine)

             X..weather.Sunshine
nobs                   366.0000
NAs                      3.0000
Minimum                  0.0000
Maximum                 13.6000
1. Quartile              5.9500
3. Quartile             10.5000
Mean                     7.9094
Median                   8.6000
Sum                   2871.1000
SE Mean                  0.1827
LCL Mean                 7.5500
UCL Mean                 8.2687
Variance                12.1210
Stdev                    3.4815
Skewness                -0.7235
Kurtosis                -0.2706
```

Explore (graphically)

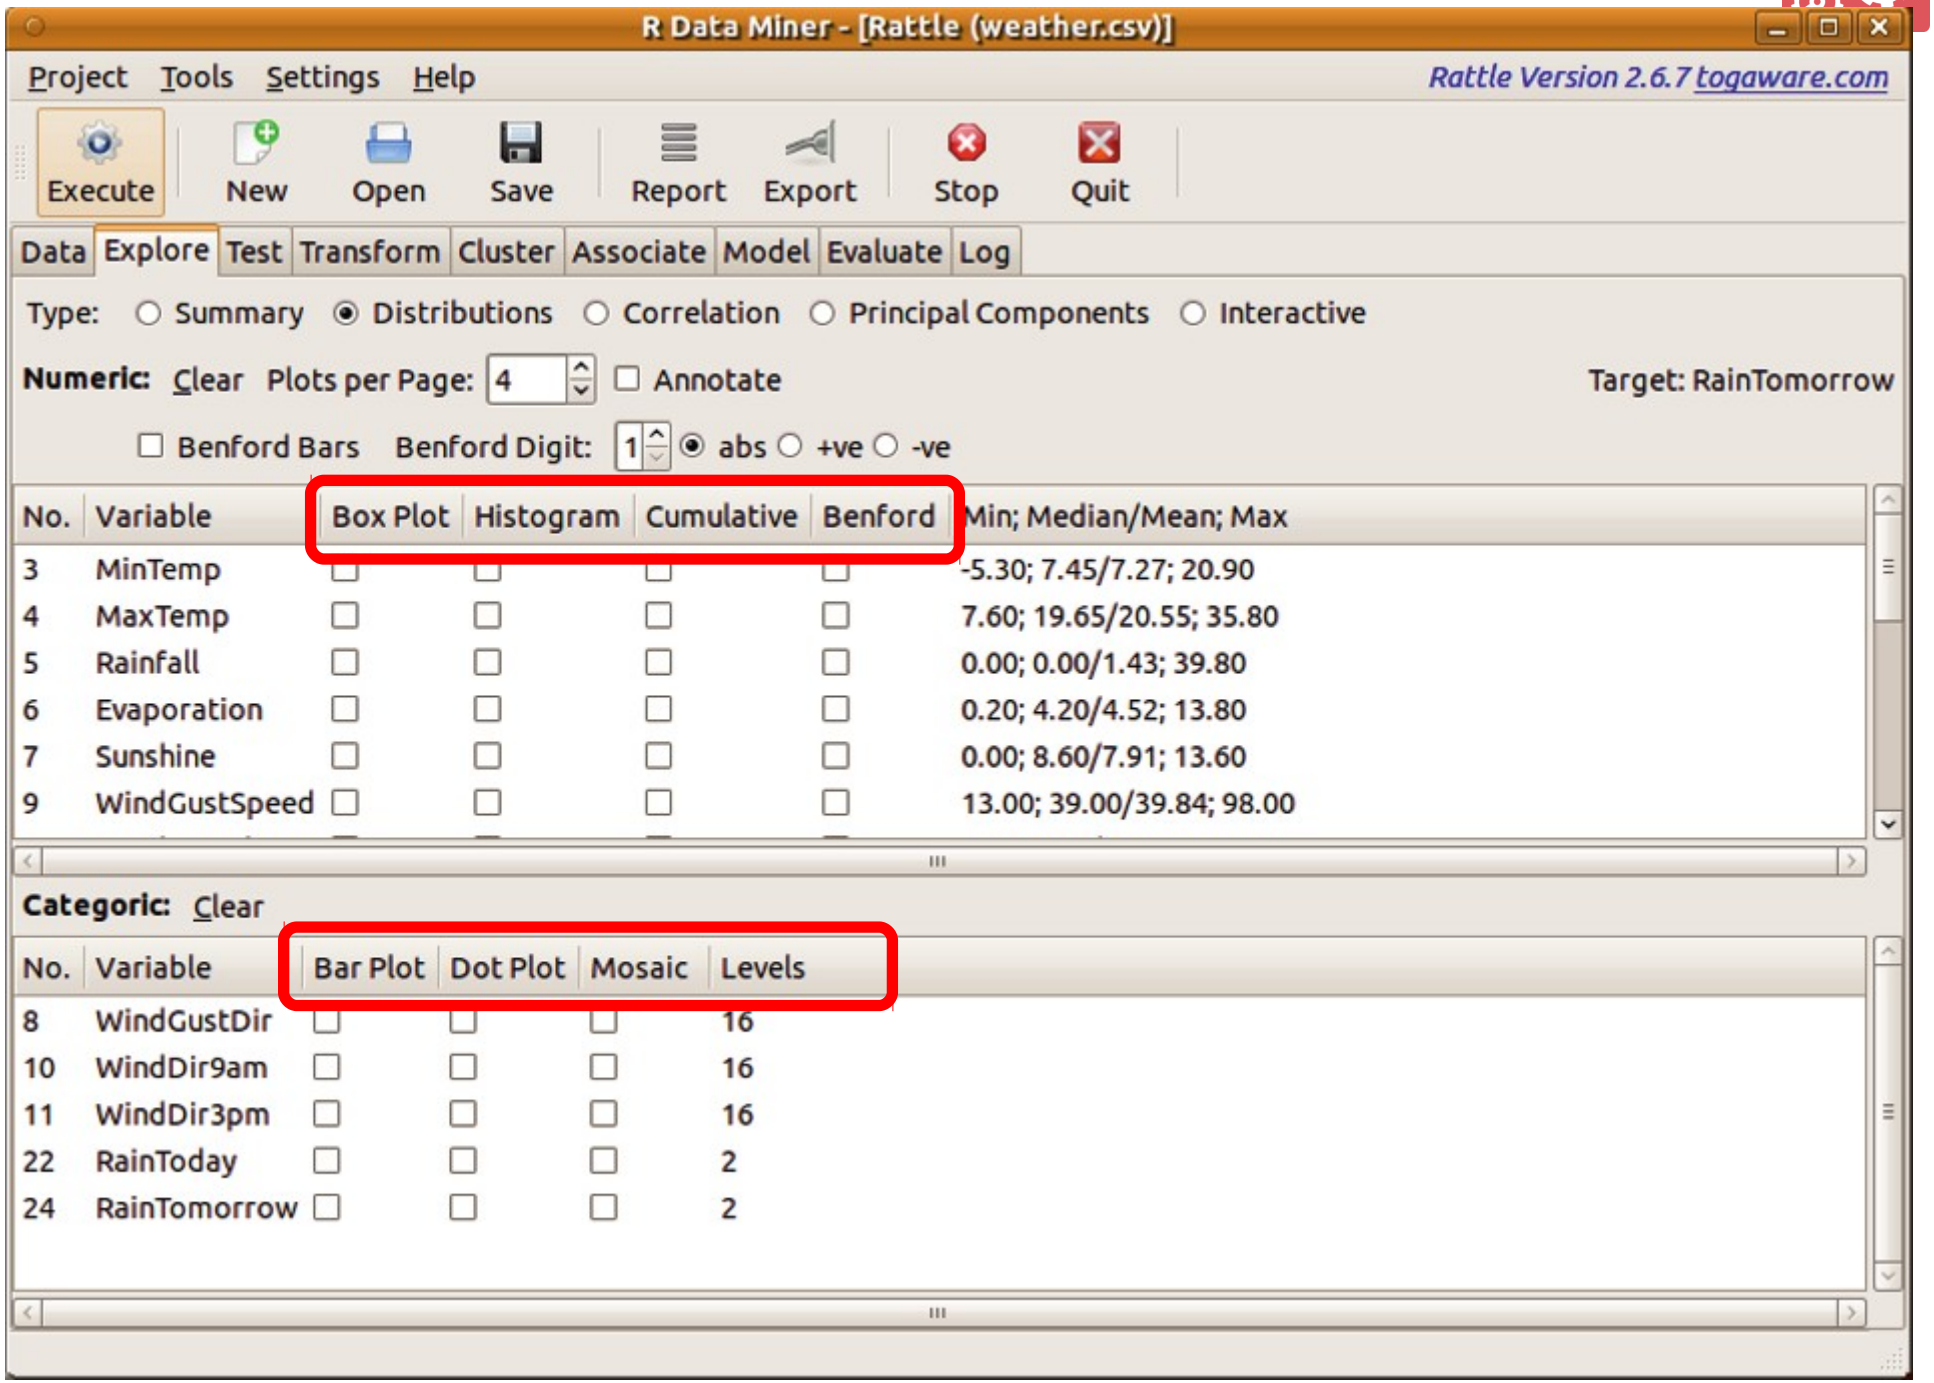

Explore (graphically)

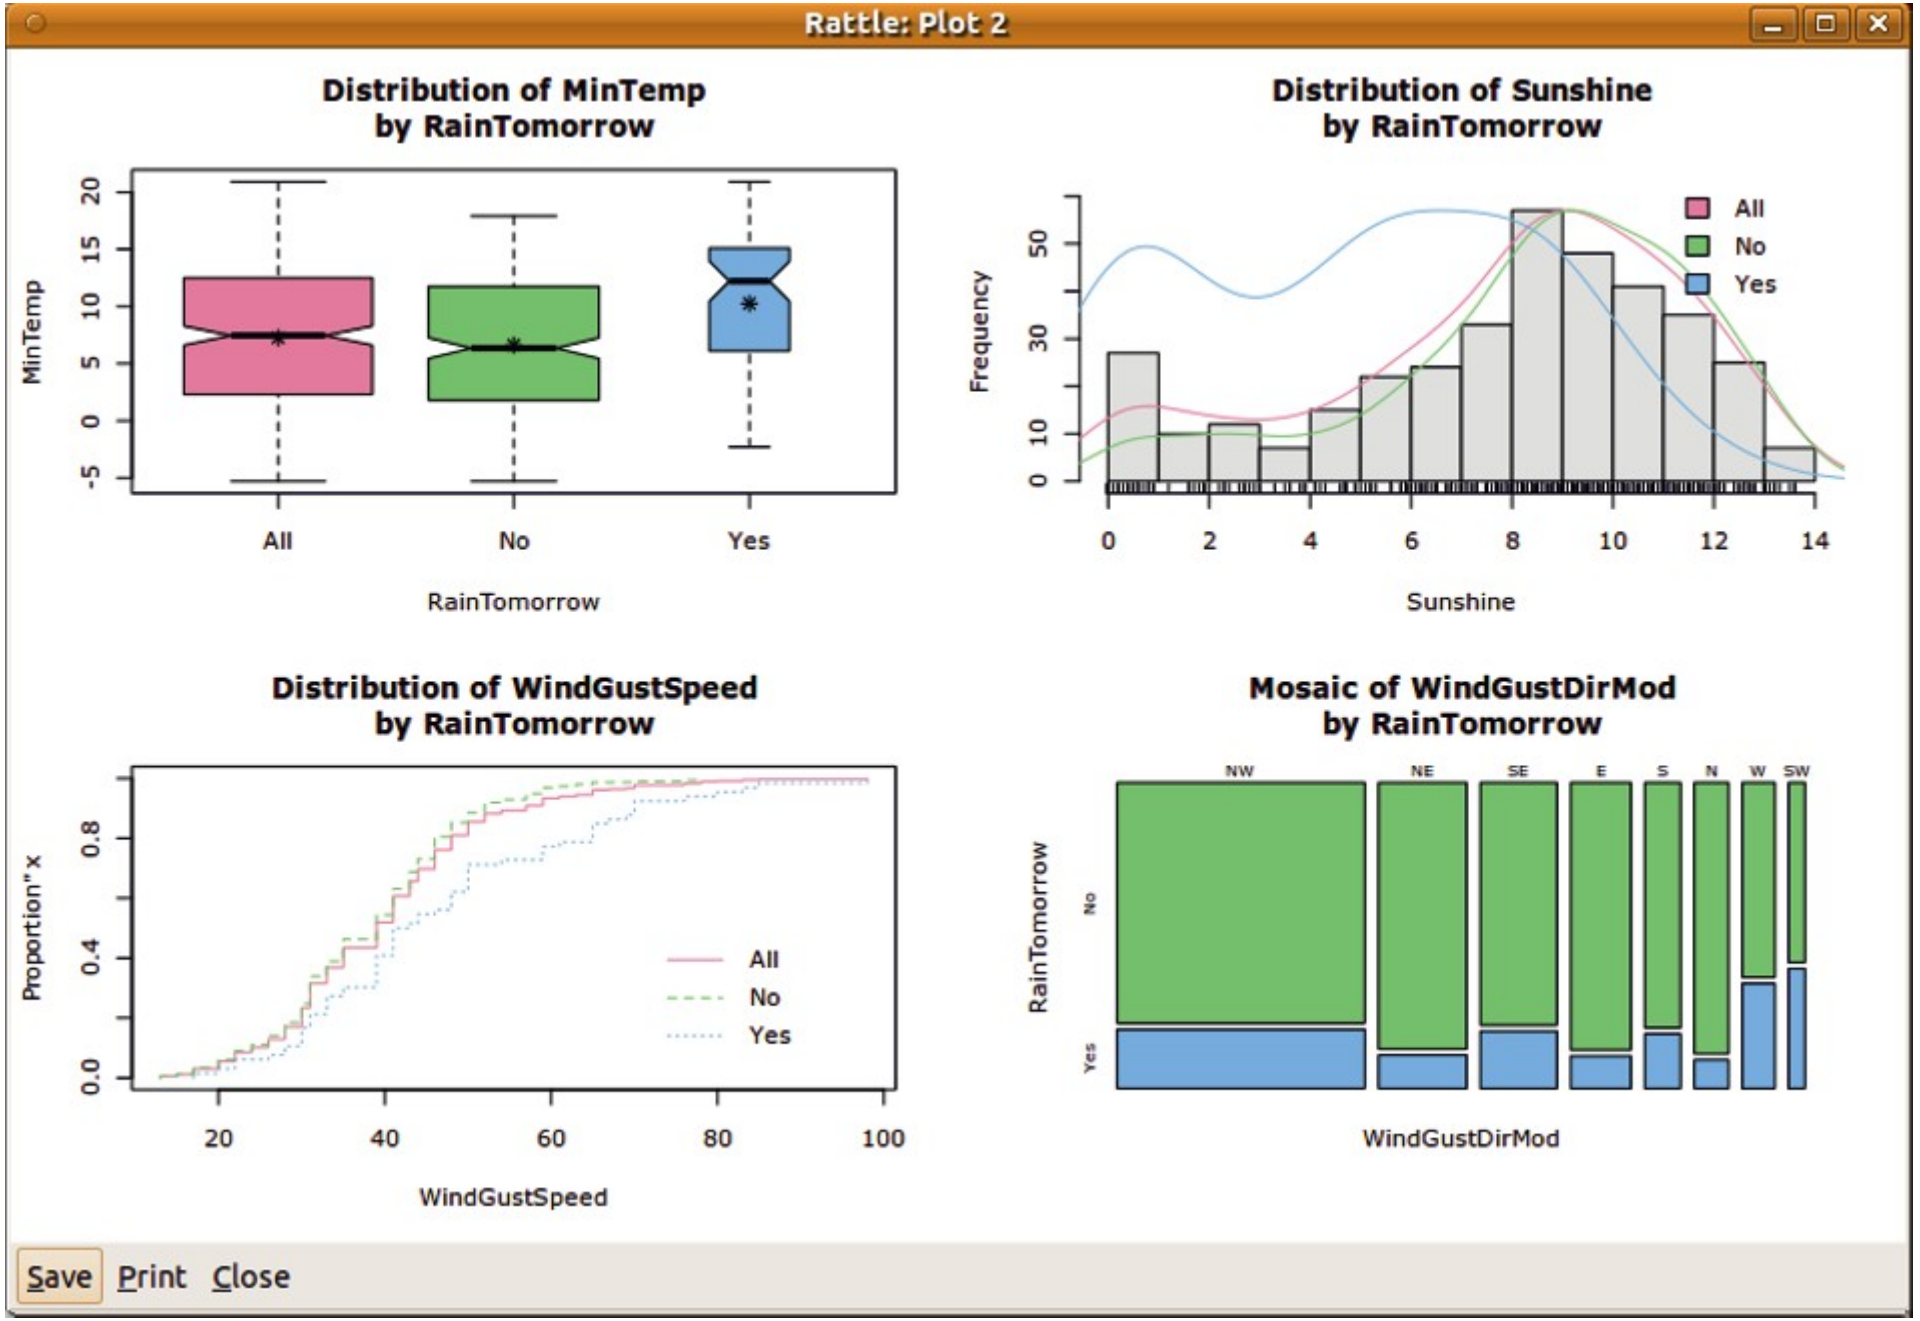

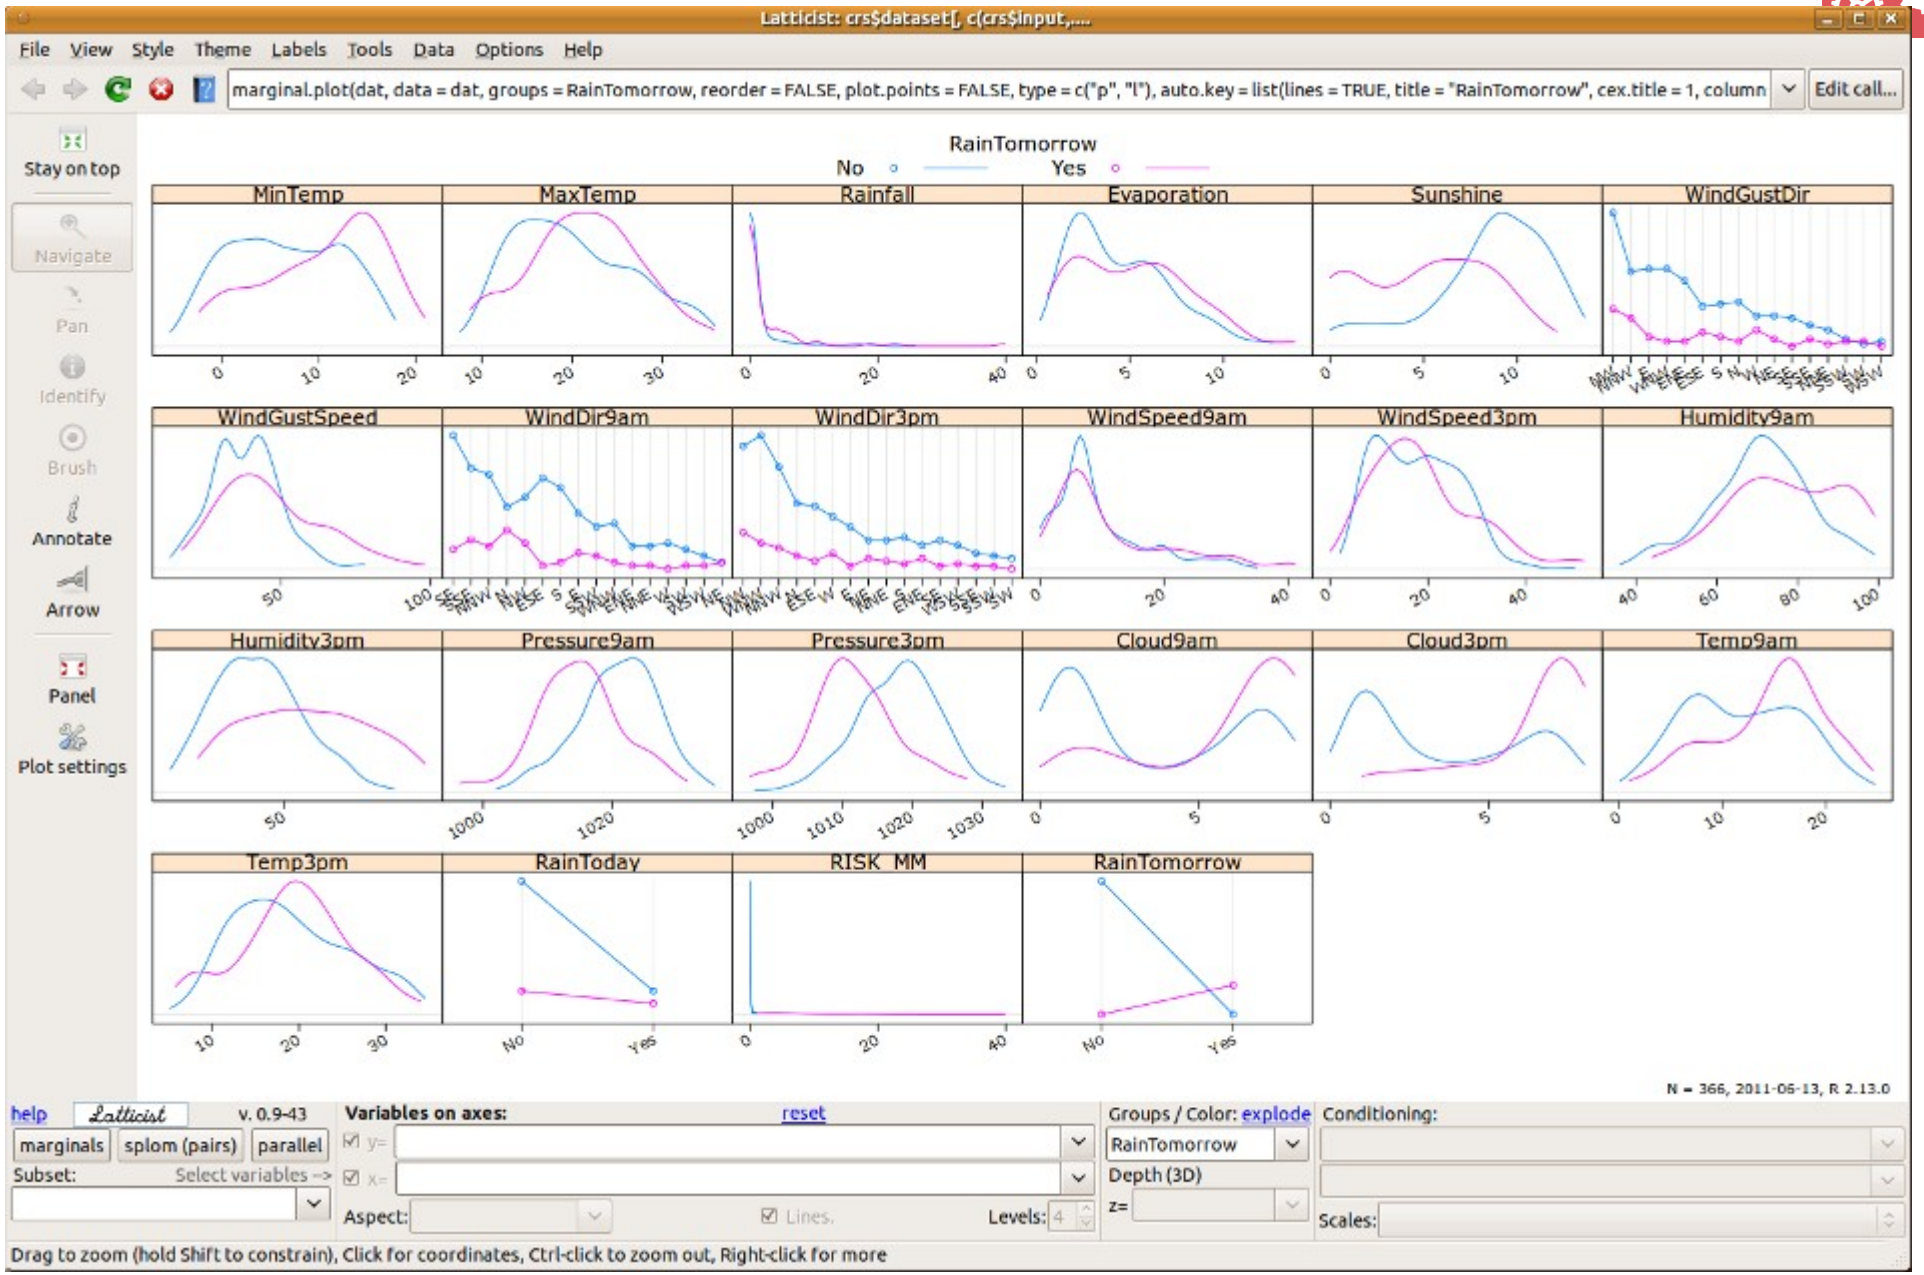

Figure 6.3: The parallel coordinates plot from latticist

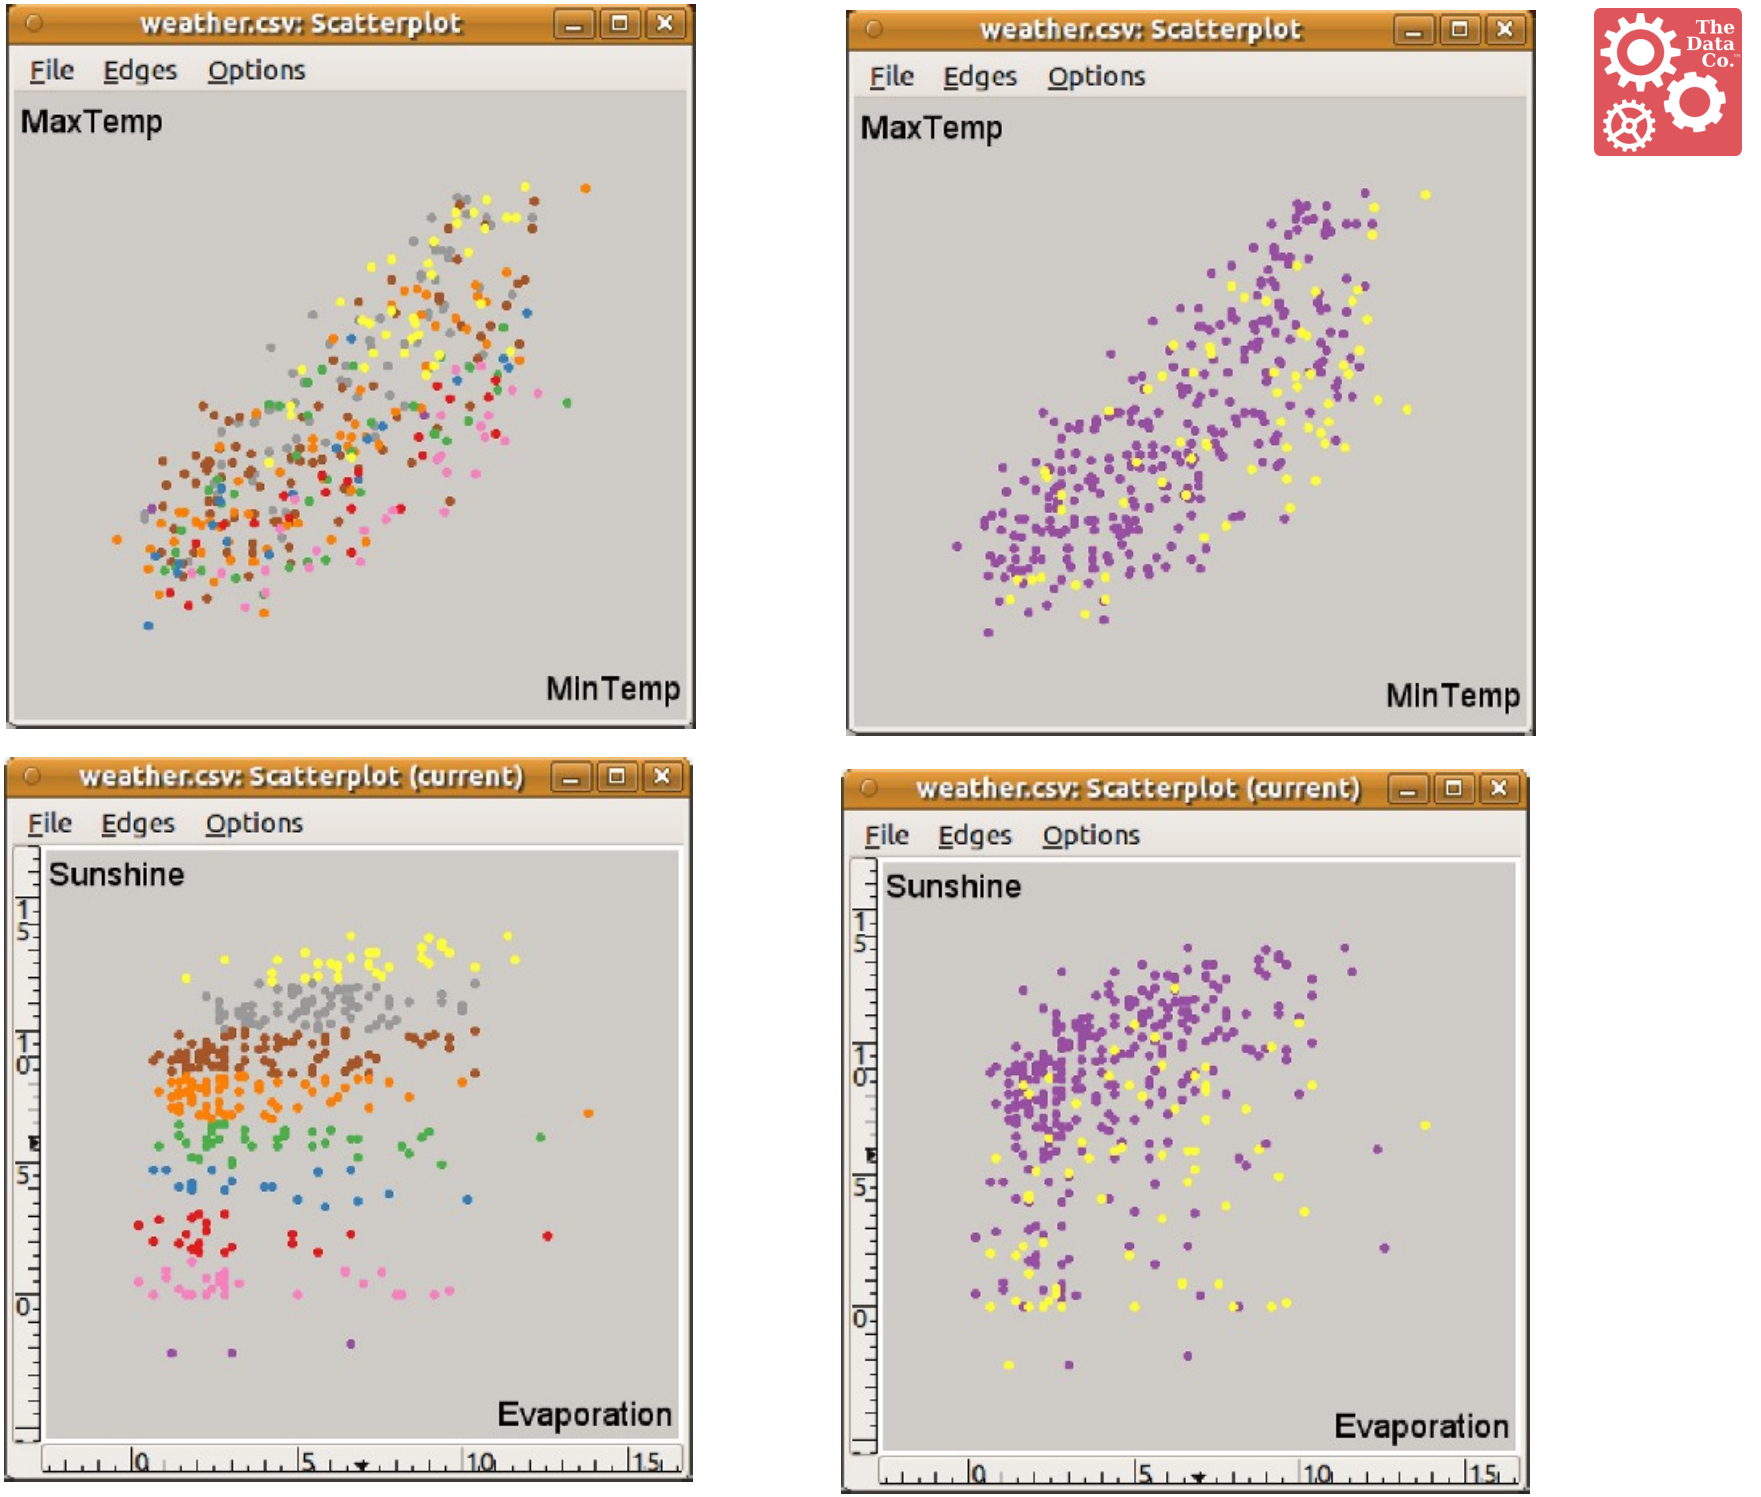

Figure 6.9: Colourful brushing of multiple scatterplots

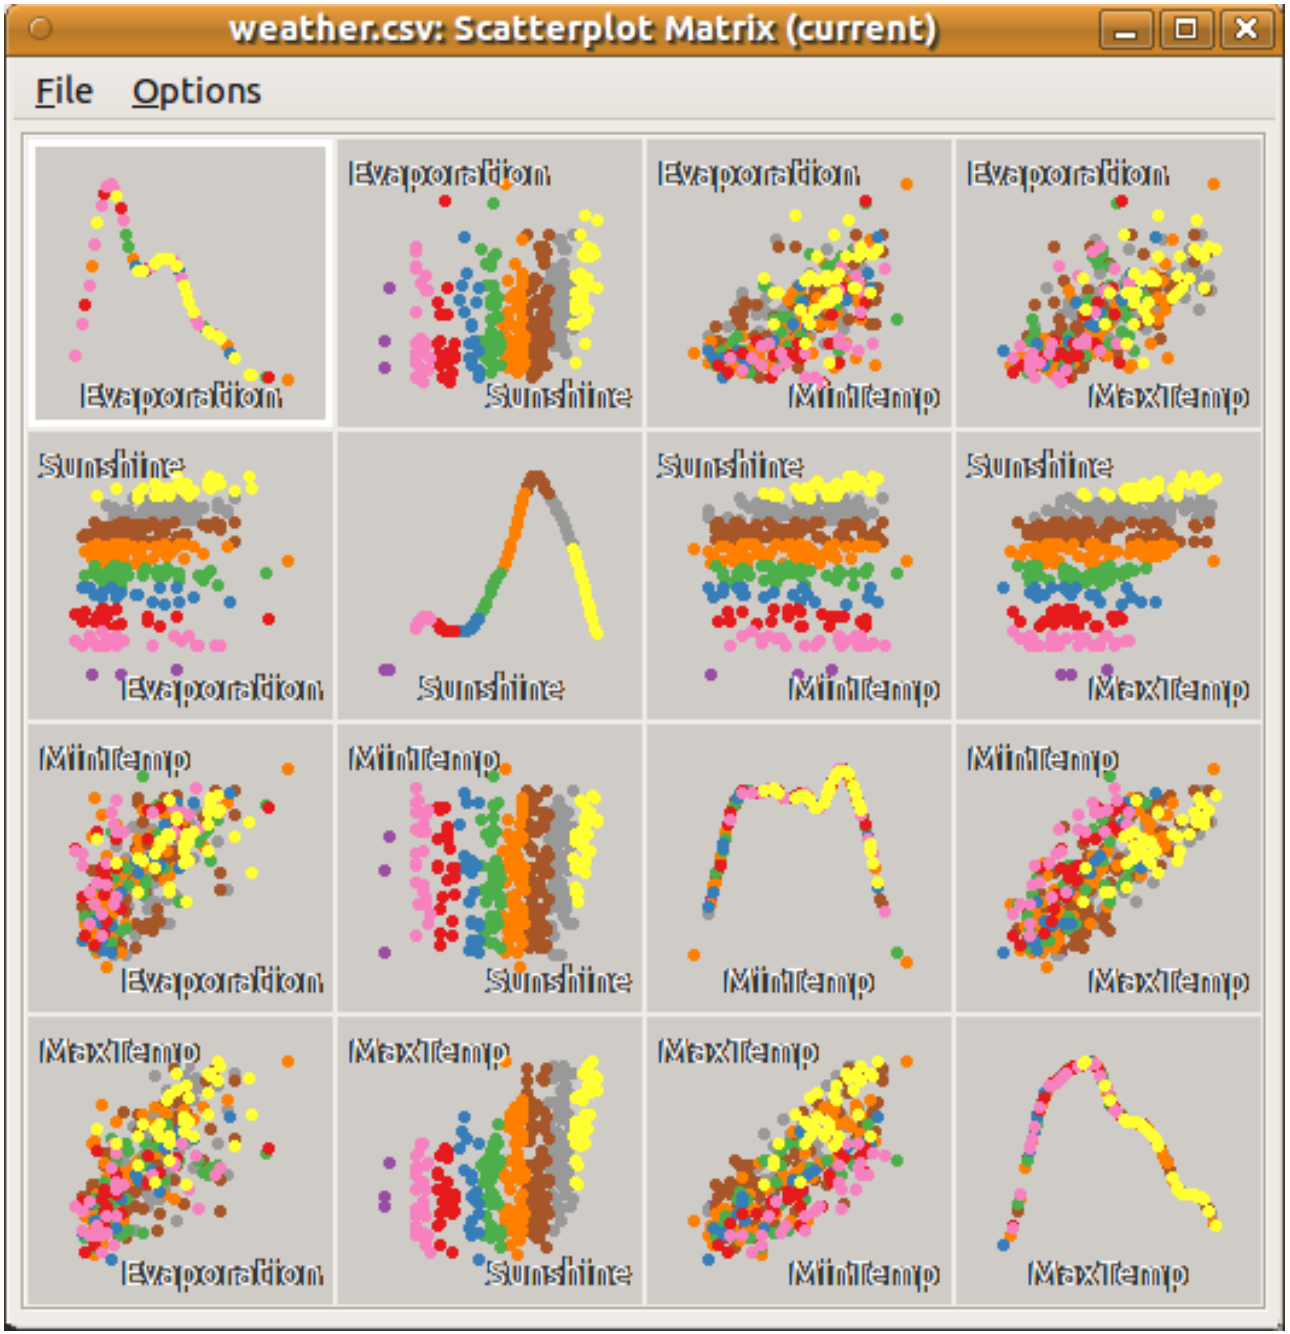

Figure 6.10: GGobi's scatter plot matrix.

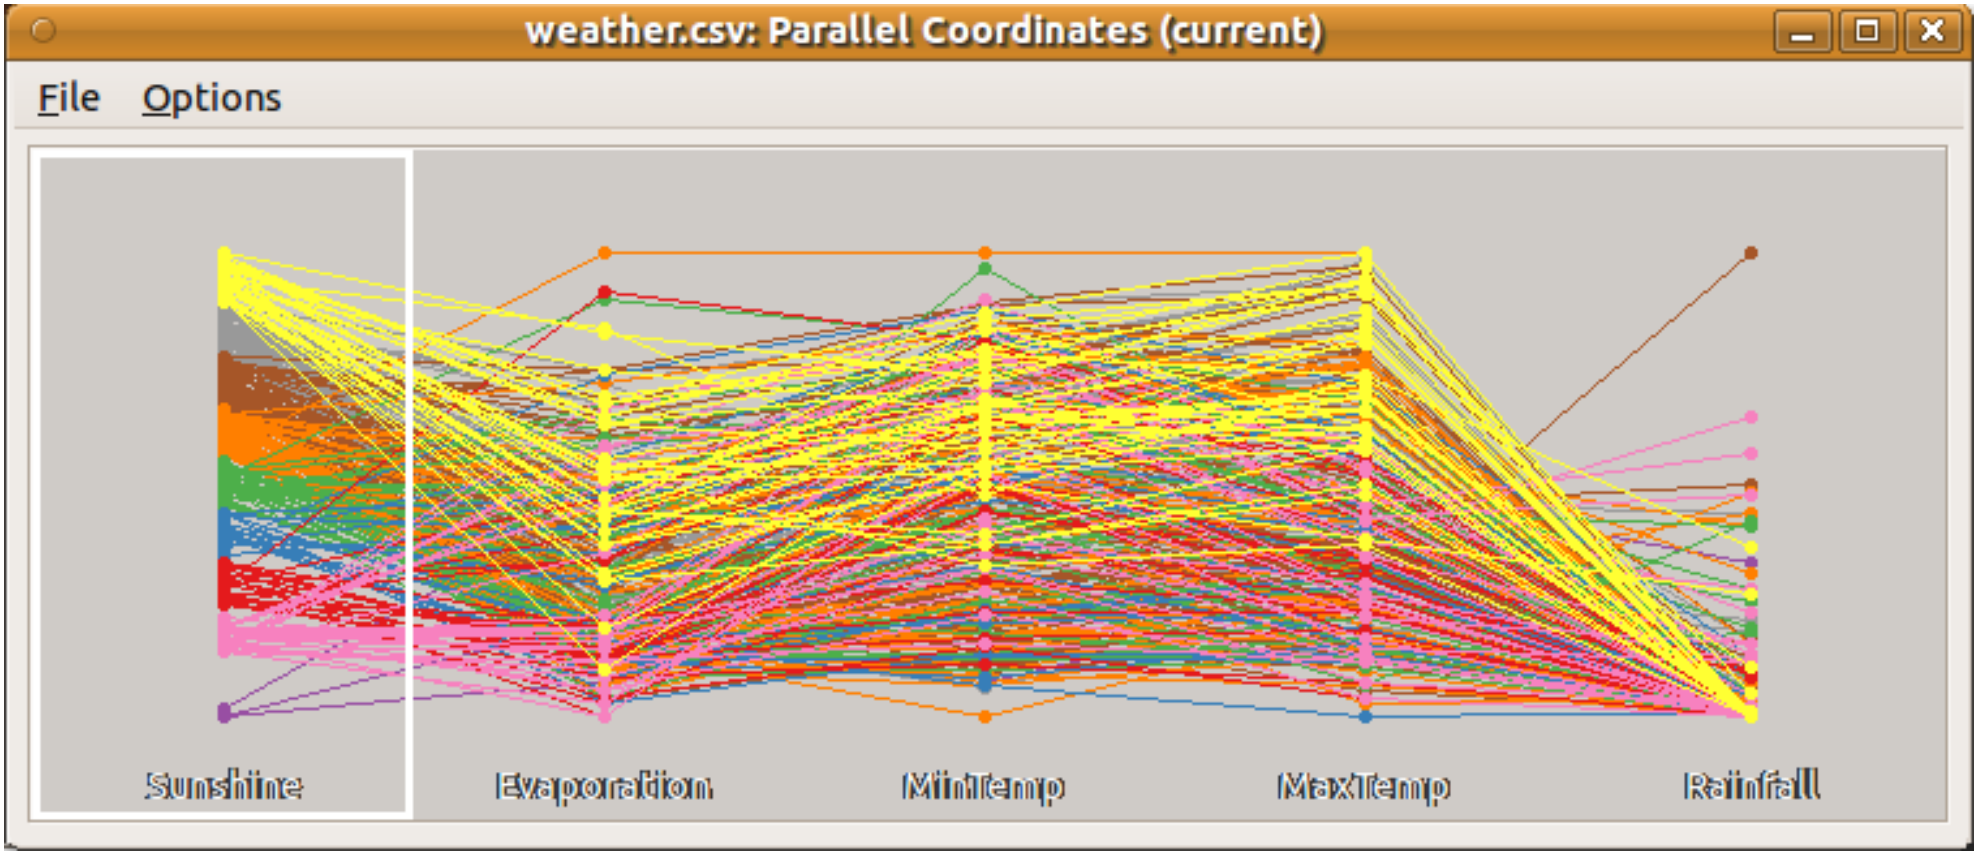

Figure 6.11: GGobi's parallel coordinates plot.

Test

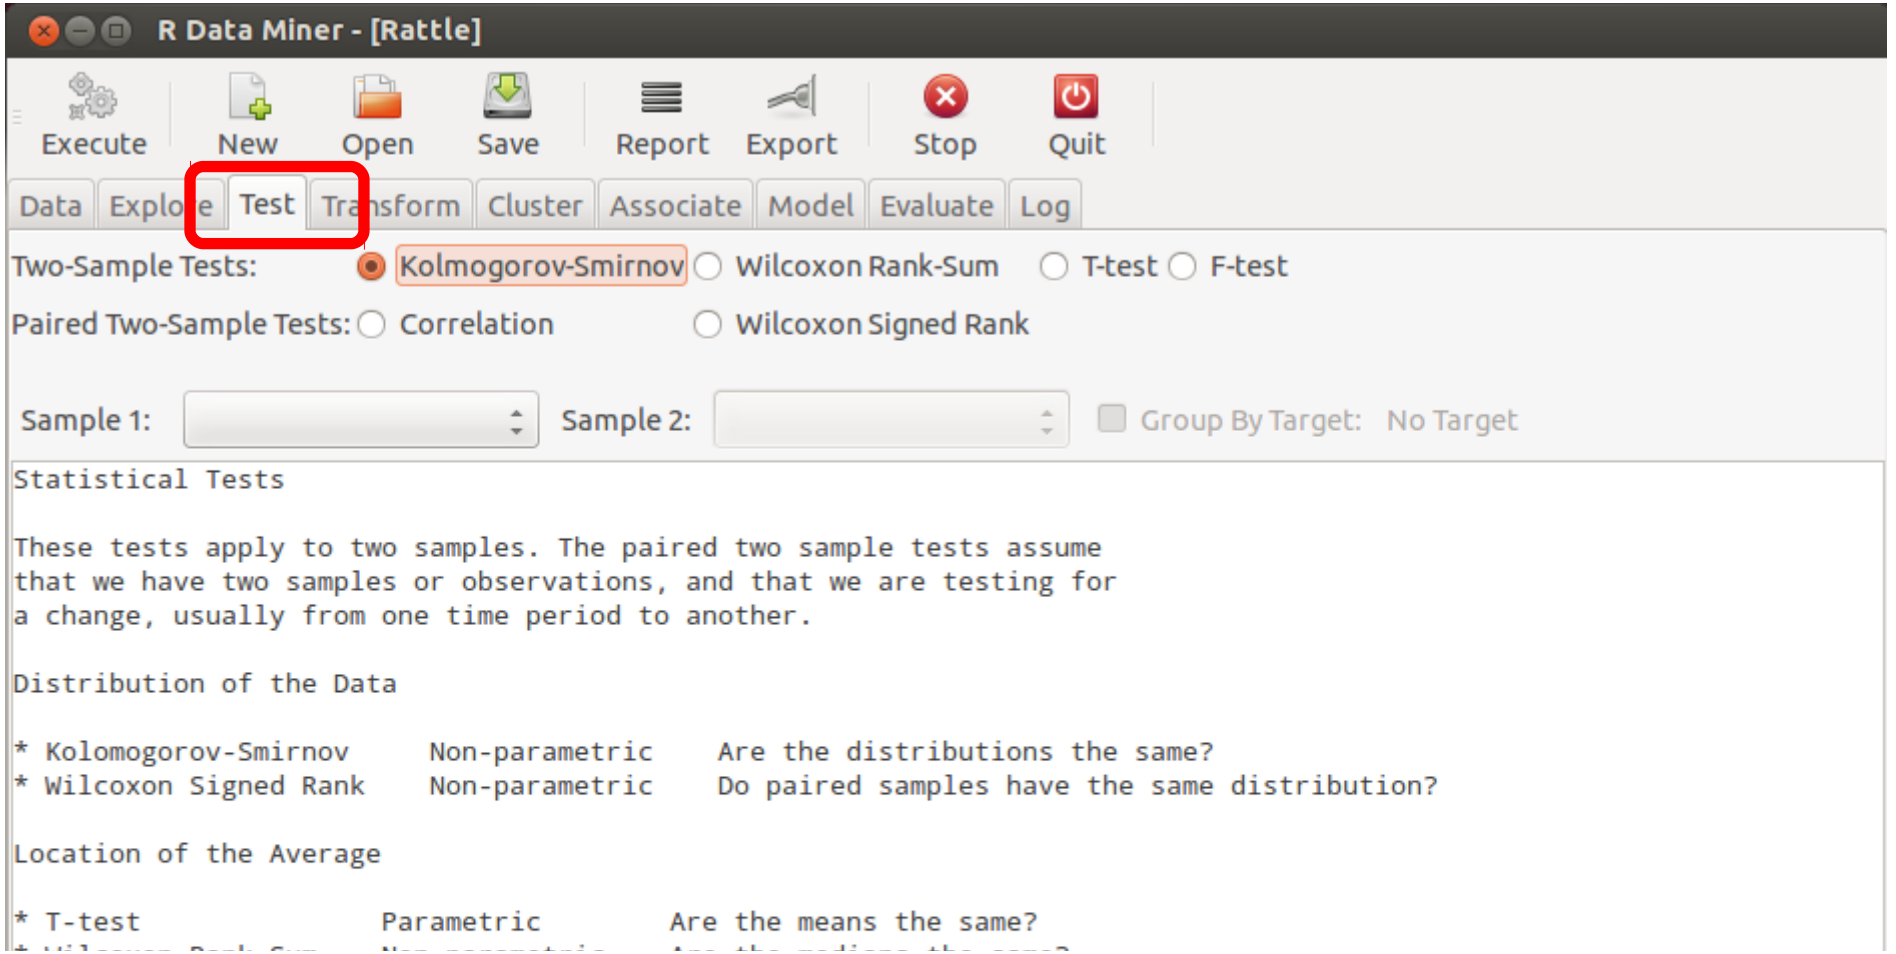

- various statistical tests, e.g. the T-test and F-test

Transform

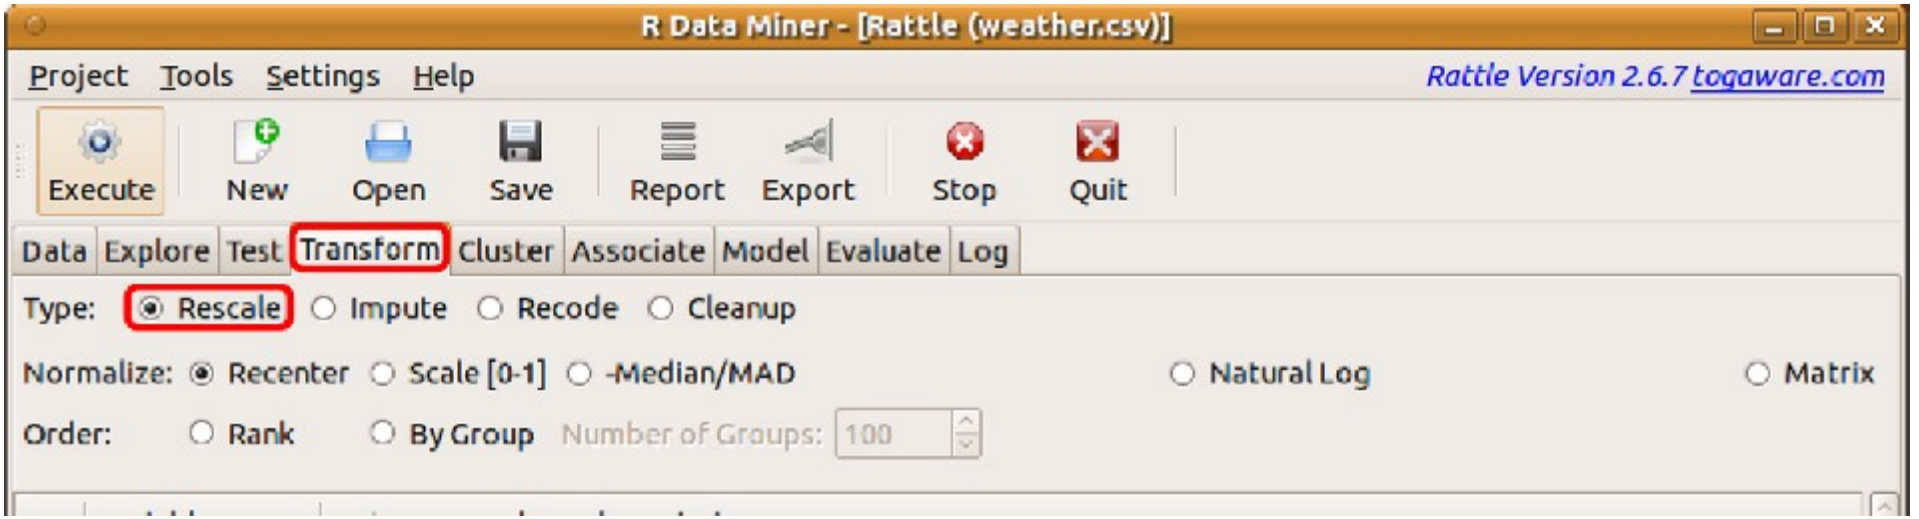

- Normalising
- Filling in missing values
- Turning numeric variables into categoric variables (and vice versa)
- Dealing with outliers
- Removing variables or observations with missing values

Cluster

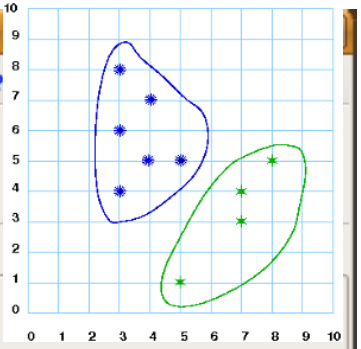

R Data Miner - [Rattle (weather.csv)]

Project Tools Settings Help Rattle Ve

Execute | New | Open | Save | Report | Export | Stop | Quit

Data | Explore | Test | Transform | Cluster | Associate | Model | Evaluate | Log

Type: KMeans Clara Hierarchical BiCluster

Number of clusters: 10 Seed: 42 Runs: 1 Stats Data Plot Discriminant Plot

Use HClust Centers Iterate Clusters

```
Cluster sizes:

[1] "23 17 22 22 17 36 23 34 22 32"

Data means:

      MinTemp       MaxTemp      Rainfall   Evaporation
        7.146        20.372         1.377         4.544
     Sunshine WindGustSpeed  WindSpeed9am  WindSpeed3pm
        8.083        39.944         9.819        18.056
  Humidity9am   Humidity3pm   Pressure9am   Pressure3pm
       71.472        43.859      1019.748      1016.979
     Cloud9am      Cloud3pm       Temp9am       Temp3pm
        3.690         3.851        12.269        19.081

Cluster centers:

   MinTemp MaxTemp Rainfall Evaporation Sunshine
1   8.5000   21.05  1.27826       6.330   10.496
2  11.6059   30.95  0.11765       7.647   11.276
3  13.4136   28.77  1.02727       6.200    9.464
4   9.1818   16.90  4.94545       3.800    2.191
5   7.7412   15.19  3.58824       3.306    5.659
```

- Allows data miners to break data into more meaningful groups and then contrast the different clusters against each other

Associate

R Data Miner - [Rattle (dvdtrans.csv)]

Project Tools Settings Help — Rattle Version 2.6.7 togaware.com

Execute | New | Open | Save | Report | Export | Stop | Quit

Data | Explore | Test | Transform | Cluster | Associate | Model | Evaluate | Log

☑ Baskets | Freq Plot | Support: 0.1000 | Confidence: 0.1000 | Show Rules

```
Summary of the Apriori Association Rules:

Number of Rules: 127

Summary of the Measures of Interestingness:

    support         confidence          lift
 Min.   :0.100   Min.   :0.100   Min.   : 0.714
 1st Qu.:0.100   1st Qu.:0.500   1st Qu.: 1.429
 Median :0.100   Median :1.000   Median : 2.500
 Mean   :0.145   Mean   :0.759   Mean   : 3.015
 3rd Qu.:0.100   3rd Qu.:1.000   3rd Qu.: 5.000
 Max.   :0.700   Max.   :1.000   Max.   :10.000

Summary of the Execution of the Apriori Command:

parameter specification:
 confidence minval smax arem  aval originalSupport
        0.1    0.1    1 none FALSE            TRUE
 support minlen maxlen target   ext
     0.1      1     10  rules FALSE

algorithmic control:
 filter tree heap memopt load sort verbose
    0.1 TRUE TRUE  FALSE TRUE    2    TRUE

apriori - find association rules with the apriori algorithm
```

The Association Rules model has been built. Time taken: 0.02 secs

Model

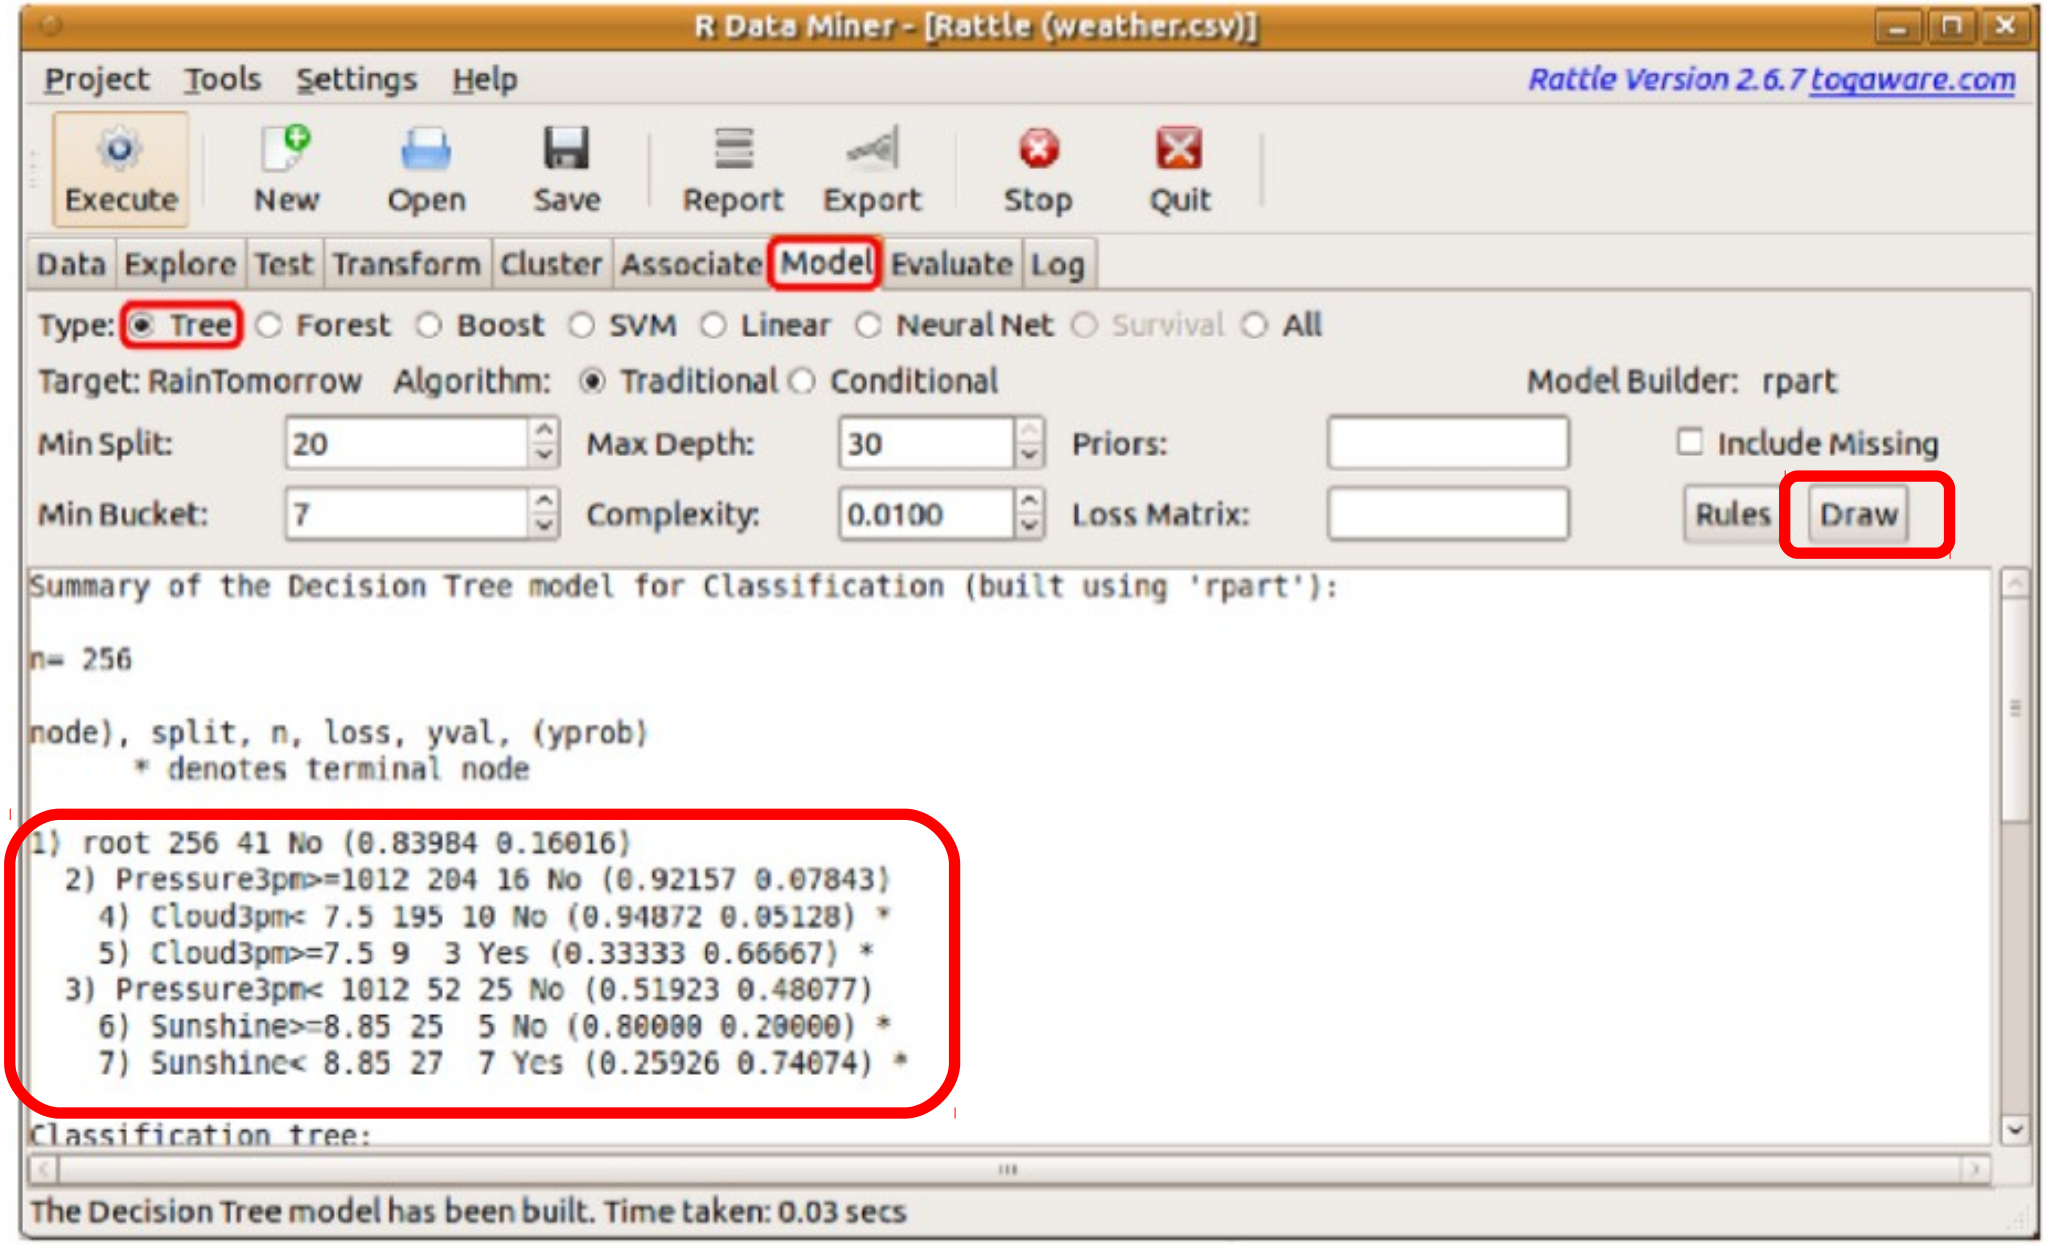
R Data Miner - [Rattle (weather.csv)]
Project Tools Settings Help
Rattle Version 2.6.7 togaware.com
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Type: Tree Forest Boost SVM Linear Neural Net Survival All
Target: RainTomorrow Algorithm: Traditional Conditional
Model Builder: rpart
Min Split: 20
Max Depth: 30
Priors:
Include Missing
Min Bucket: 7
Complexity: 0.0100
Loss Matrix:
Rules Draw
Summary of the Decision Tree model for Classification (built using 'rpart'):
n= 256
node), split, n, loss, yval, (yprob)
* denotes terminal node
1) root 256 41 No (0.83984 0.16016)
2) Pressure3pm>=1012 204 16 No (0.92157 0.07843)
4) Cloud3pm< 7.5 195 10 No (0.94872 0.05128) *
5) Cloud3pm>=7.5 9 3 Yes (0.33333 0.66667) *
3) Pressure3pm< 1012 52 25 No (0.51923 0.48077)
6) Sunshine>=8.85 25 5 No (0.80000 0.20000) *
7) Sunshine< 8.85 27 7 Yes (0.25926 0.74074) *
Classification tree:
The Decision Tree model has been built. Time taken: 0.03 secs

Decision Trees

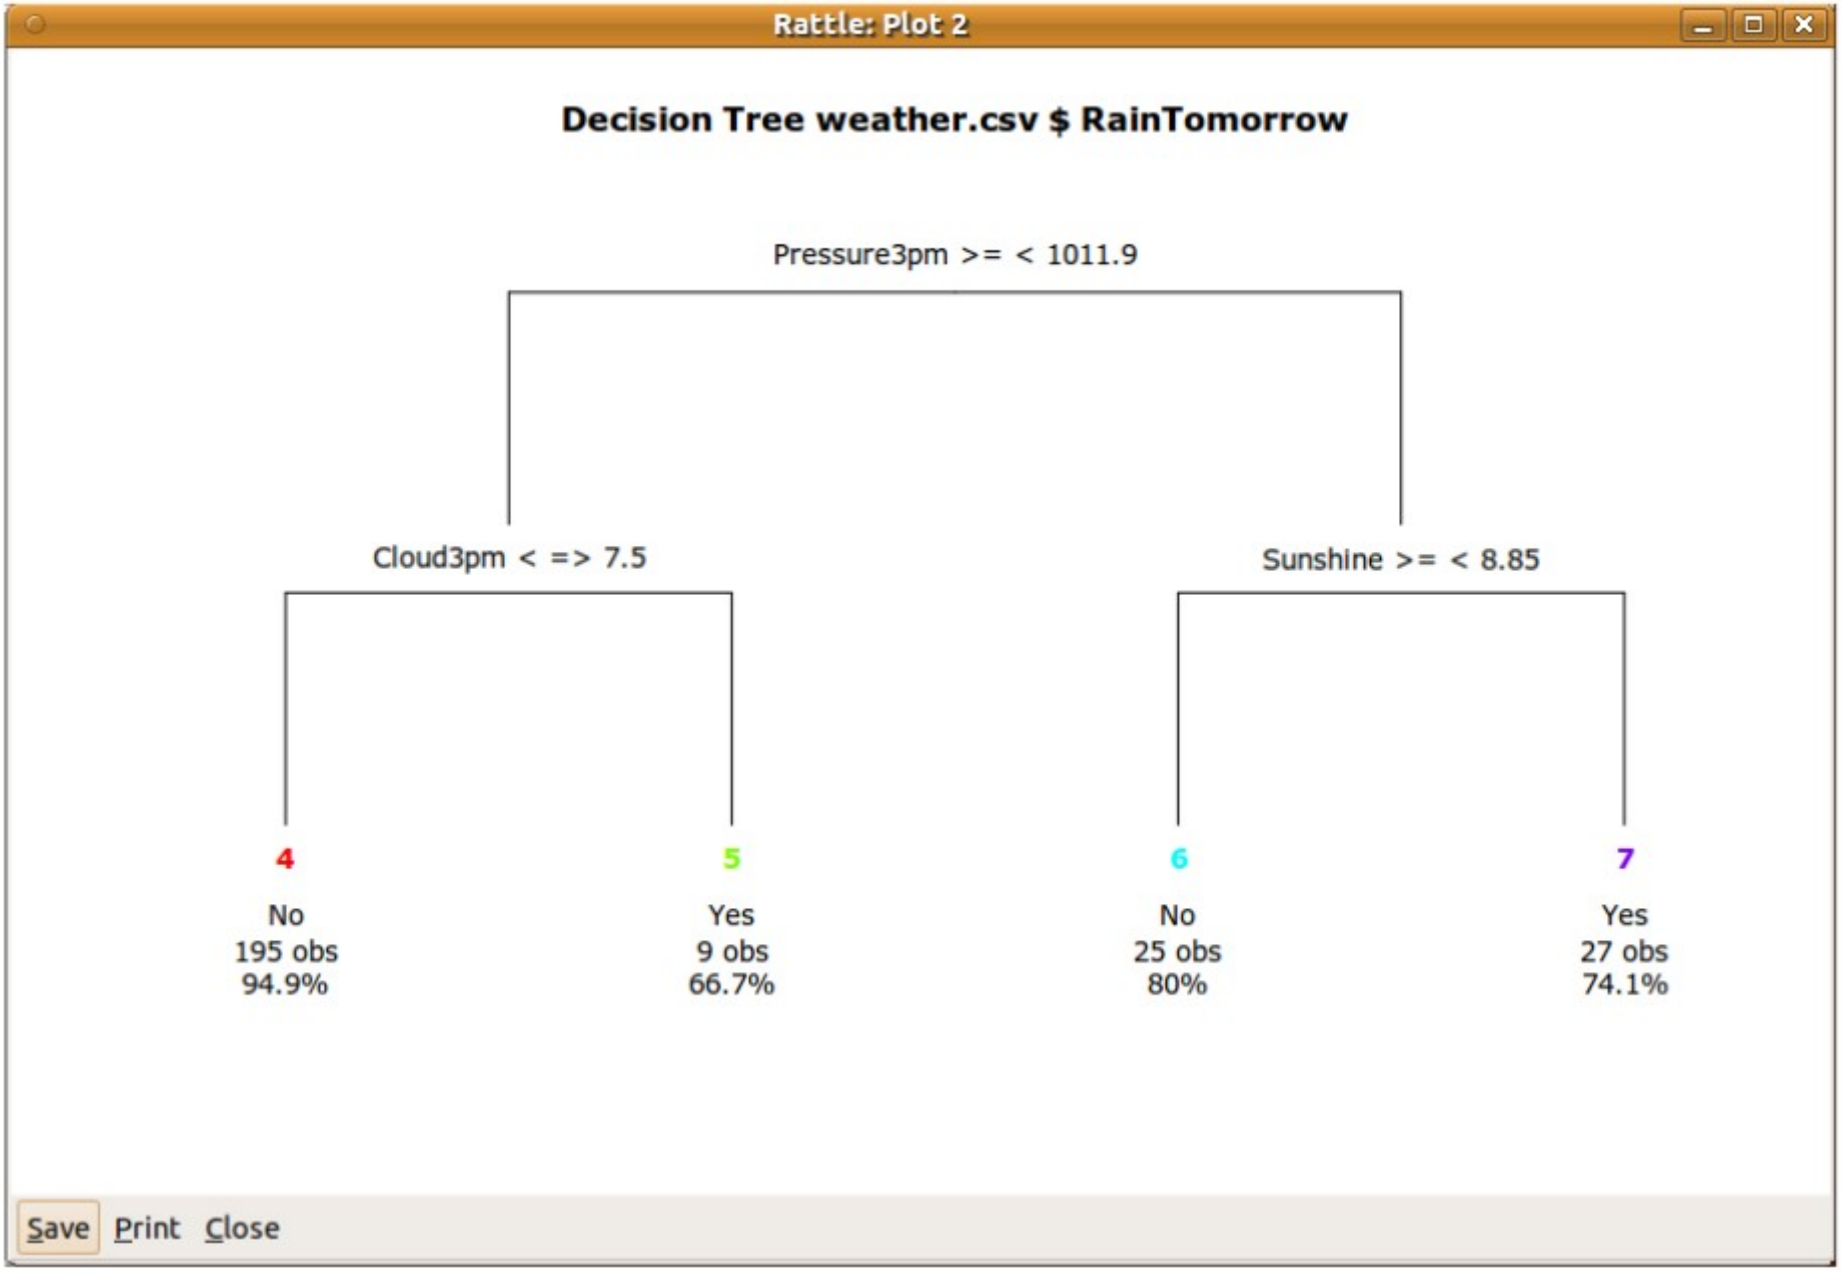

Boosting

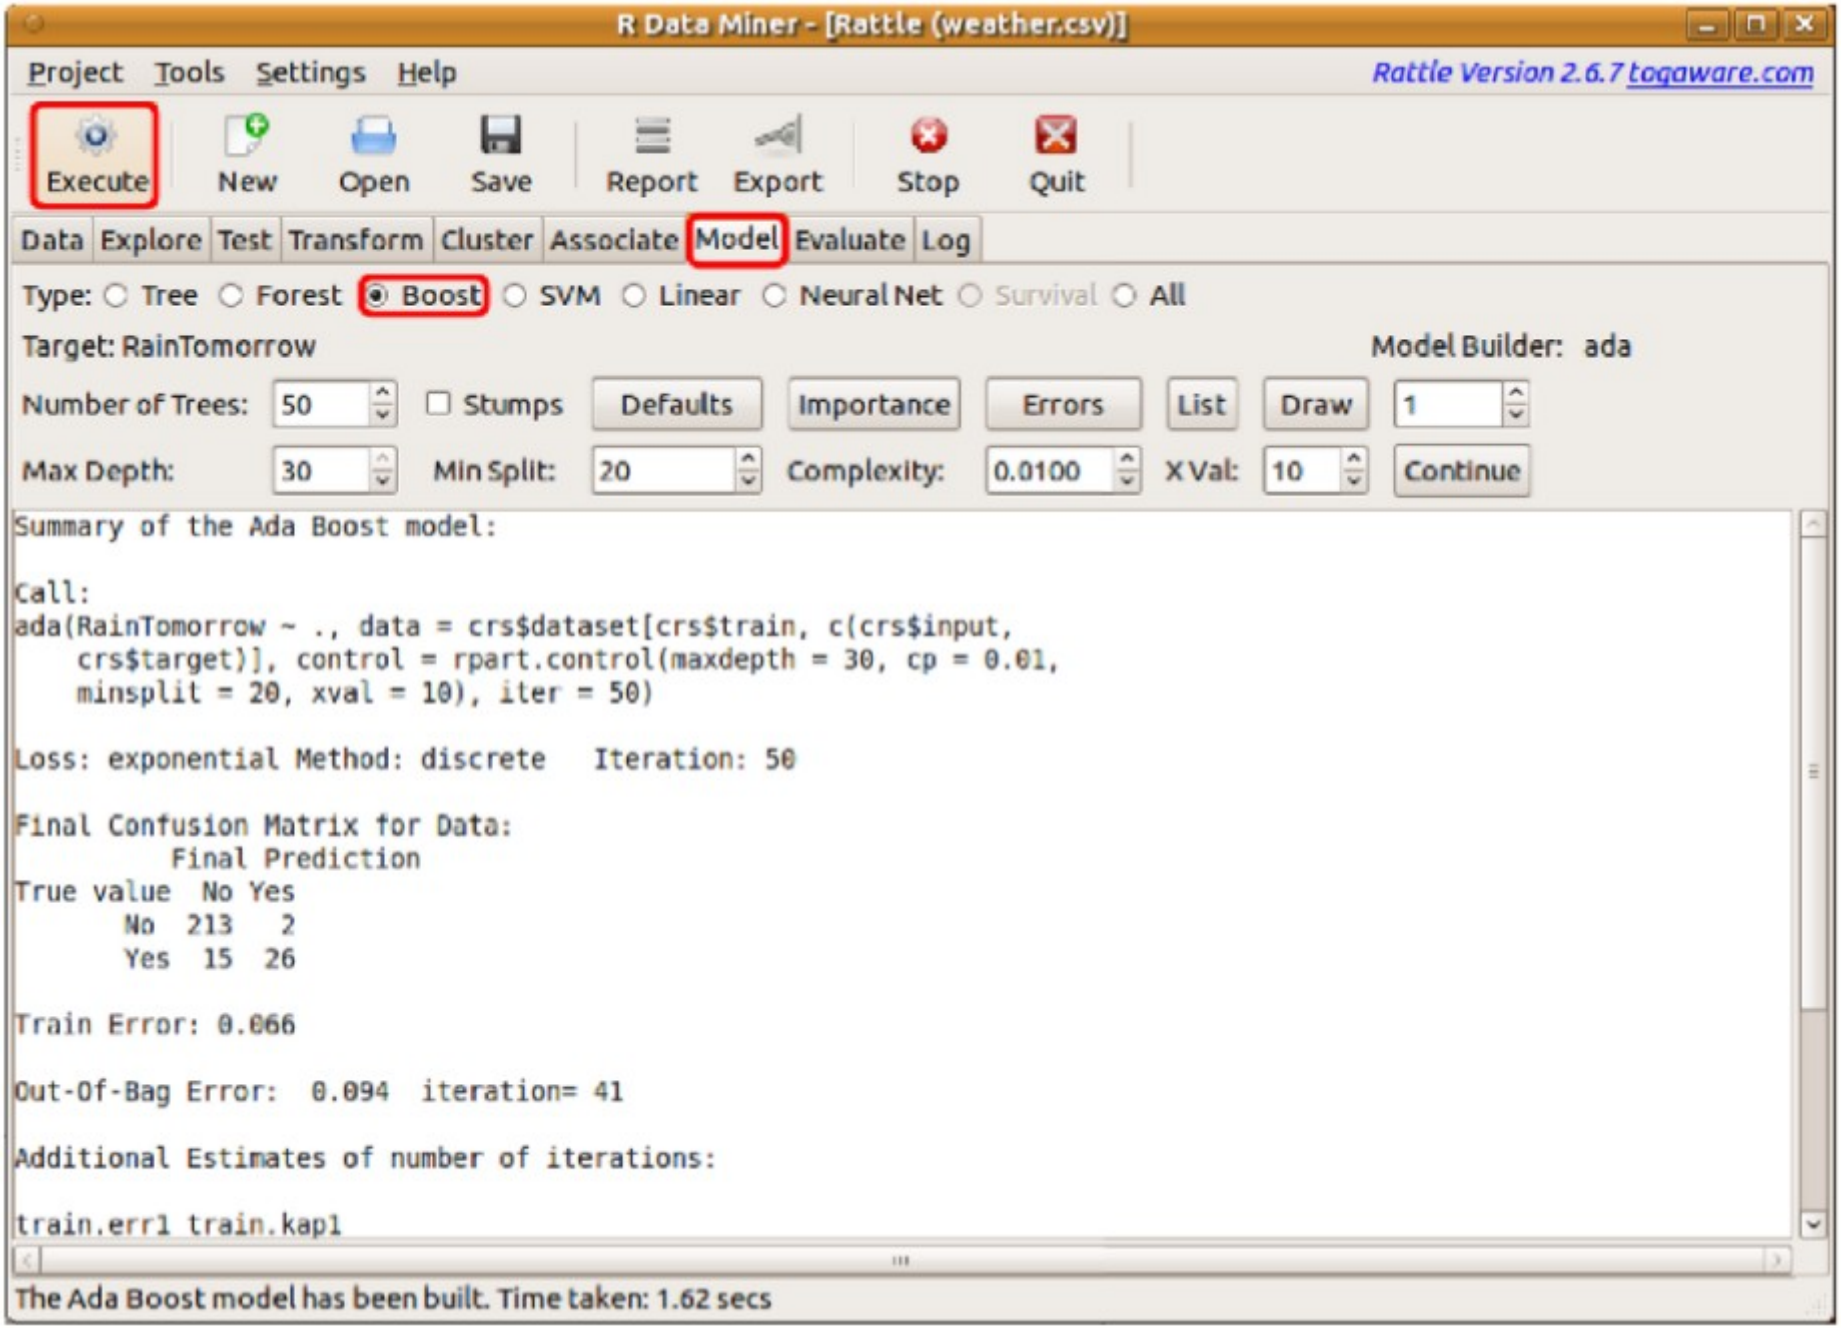
R Data Miner - [Rattle (weather.csv)]
Project Tools Settings Help
Rattle Version 2.6.7 togaware.com
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Type: Tree Forest Boost SVM Linear Neural Net Survival All
Target: RainTomorrow
Model Builder: ada
Number of Trees: 50 Stumps Defaults Importance Errors List Draw 1
Max Depth: 30 Min Split: 20 Complexity: 0.0100 X Val: 10 Continue
Summary of the Ada Boost model:
Call:
ada(RainTomorrow ~ ., data = crs$dataset[crs$train, c(crs$input,
crs$target)], control = rpart.control(maxdepth = 30, cp = 0.01,
minsplit = 20, xval = 10), iter = 50)
Loss: exponential Method: discrete Iteration: 50
Final Confusion Matrix for Data:
Final Prediction
True value No Yes
No 213 2
Yes 15 26
Train Error: 0.066
Out-Of-Bag Error: 0.094 iteration= 41
Additional Estimates of number of iterations:
train.err1 train.kap1
The Ada Boost model has been built. Time taken: 1.62 secs

Boosting

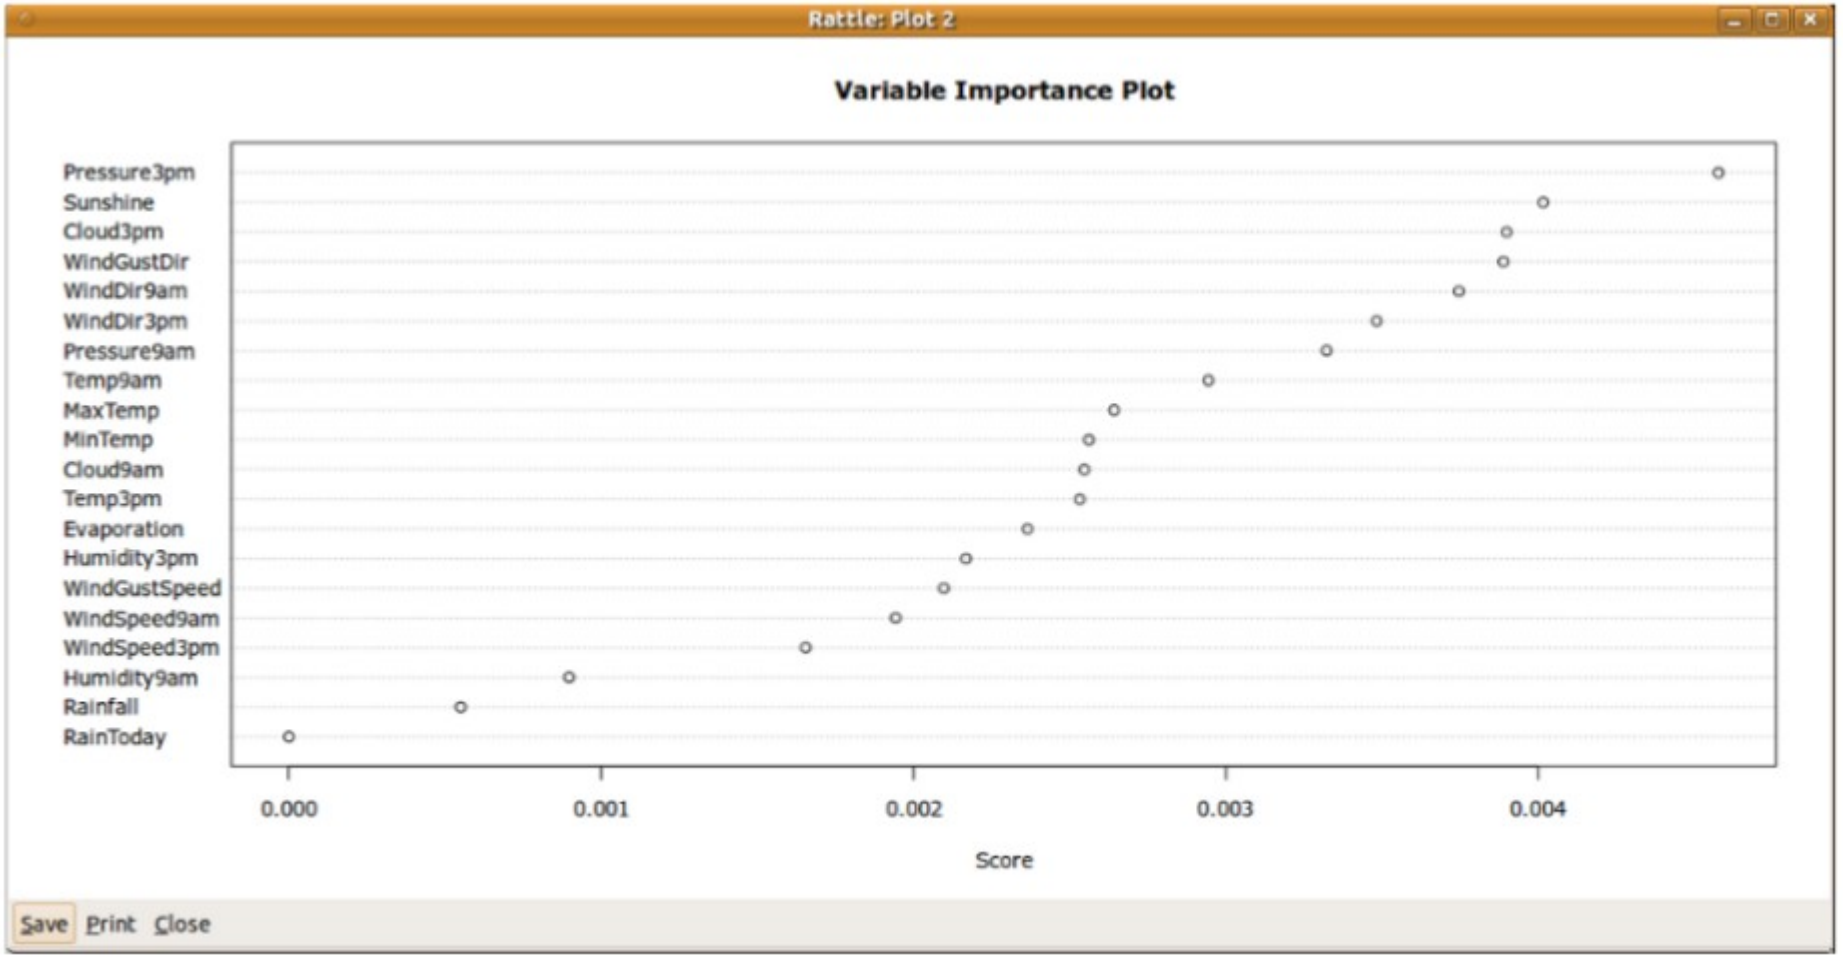

Figure 13.3: The variable importance plot for a boosted model.

Evaluate

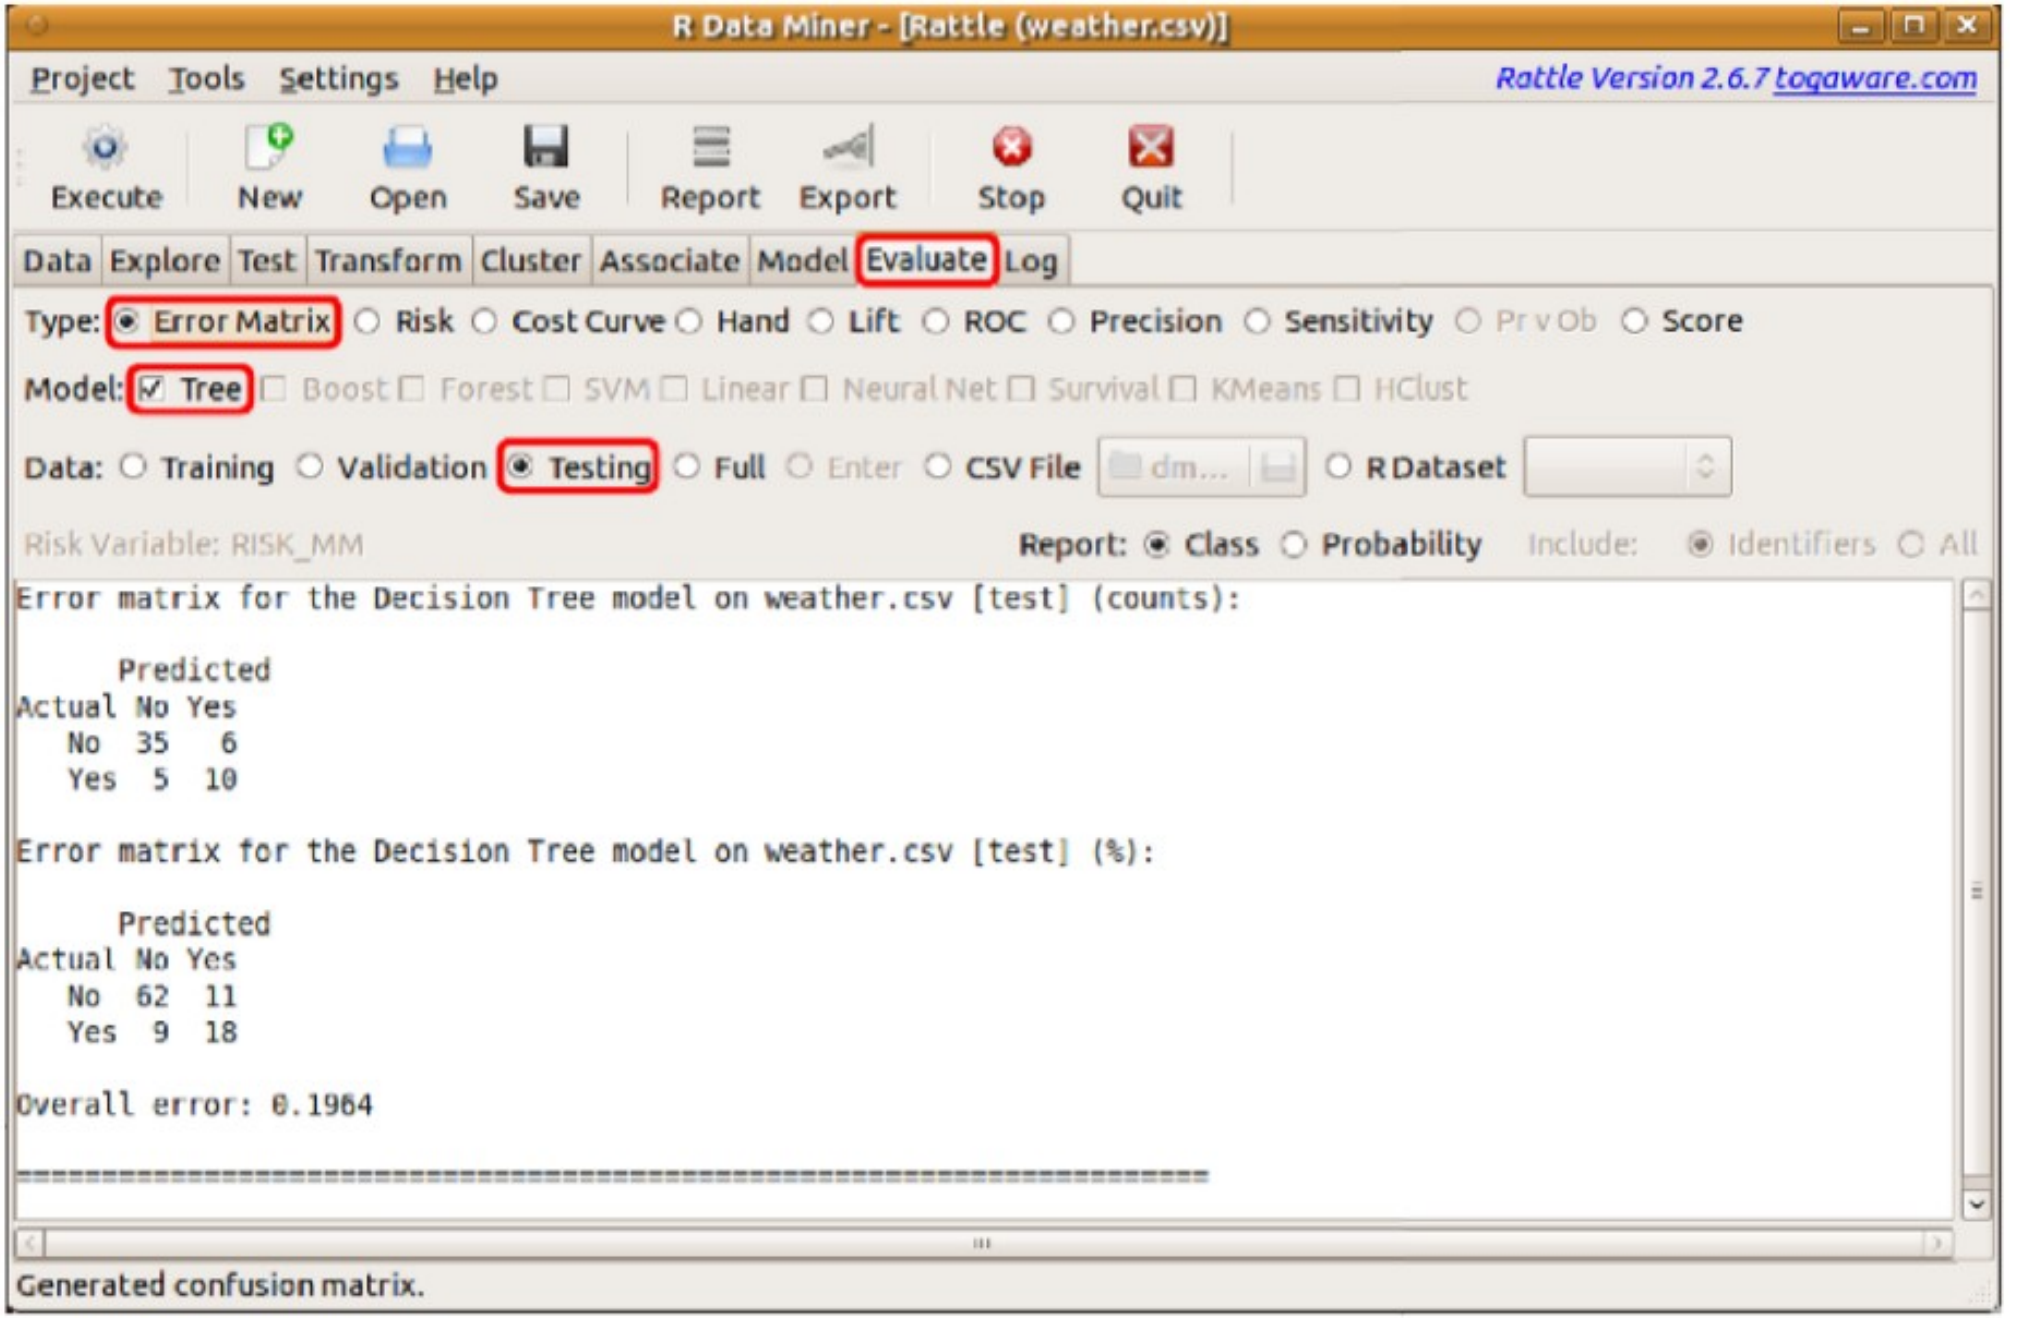
R Data Miner - [Rattle (weather.csv)]
Project Tools Settings Help
Rattle Version 2.6.7 togaware.com
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Type: Error Matrix Risk Cost Curve Hand Lift ROC Precision Sensitivity Pr v Ob Score
Model: Tree Boost Forest SVM Linear Neural Net Survival KMeans HClust
Data: Training Validation Testing Full Enter CSV File dm... R Dataset
Risk Variable: RISK_MM
Report: Class Probability
Include: Identifiers All
Error matrix for the Decision Tree model on weather.csv [test] (counts):
Predicted
Actual No Yes
No 35 6
Yes 5 10
Error matrix for the Decision Tree model on weather.csv [test] (%):
Predicted
Actual No Yes
No 62 11
Yes 9 18
Overall error: 0.1964
Generated confusion matrix.

Evaluation Using the Training Dataset:

Count		Predict	
		No	Yes
Actual	No	205	10
	Yes	15	26

Percentage		Predict	
		No	Yes
Actual	No	80	4
	Yes	6	10

Evaluation Using the Validation Dataset:

Count		Predict	
		No	Yes
Actual	No	39	5
	Yes	5	5

Percentage		Predict	
		No	Yes
Actual	No	72	9
	Yes	9	9

Evaluation Using the Testing Dataset:

Count		Predict	
		No	Yes
Actual	No	35	6
	Yes	5	10

Percentage		Predict	
		No	Yes
Actual	No	62	11
	Yes	9	18

Evaluation Using the Full Dataset:

Count		Predict	
		No	Yes
Actual	No	279	21
	Yes	25	41

Percentage		Predict	
		No	Yes
Actual	No	76	6
	Yes	7	11

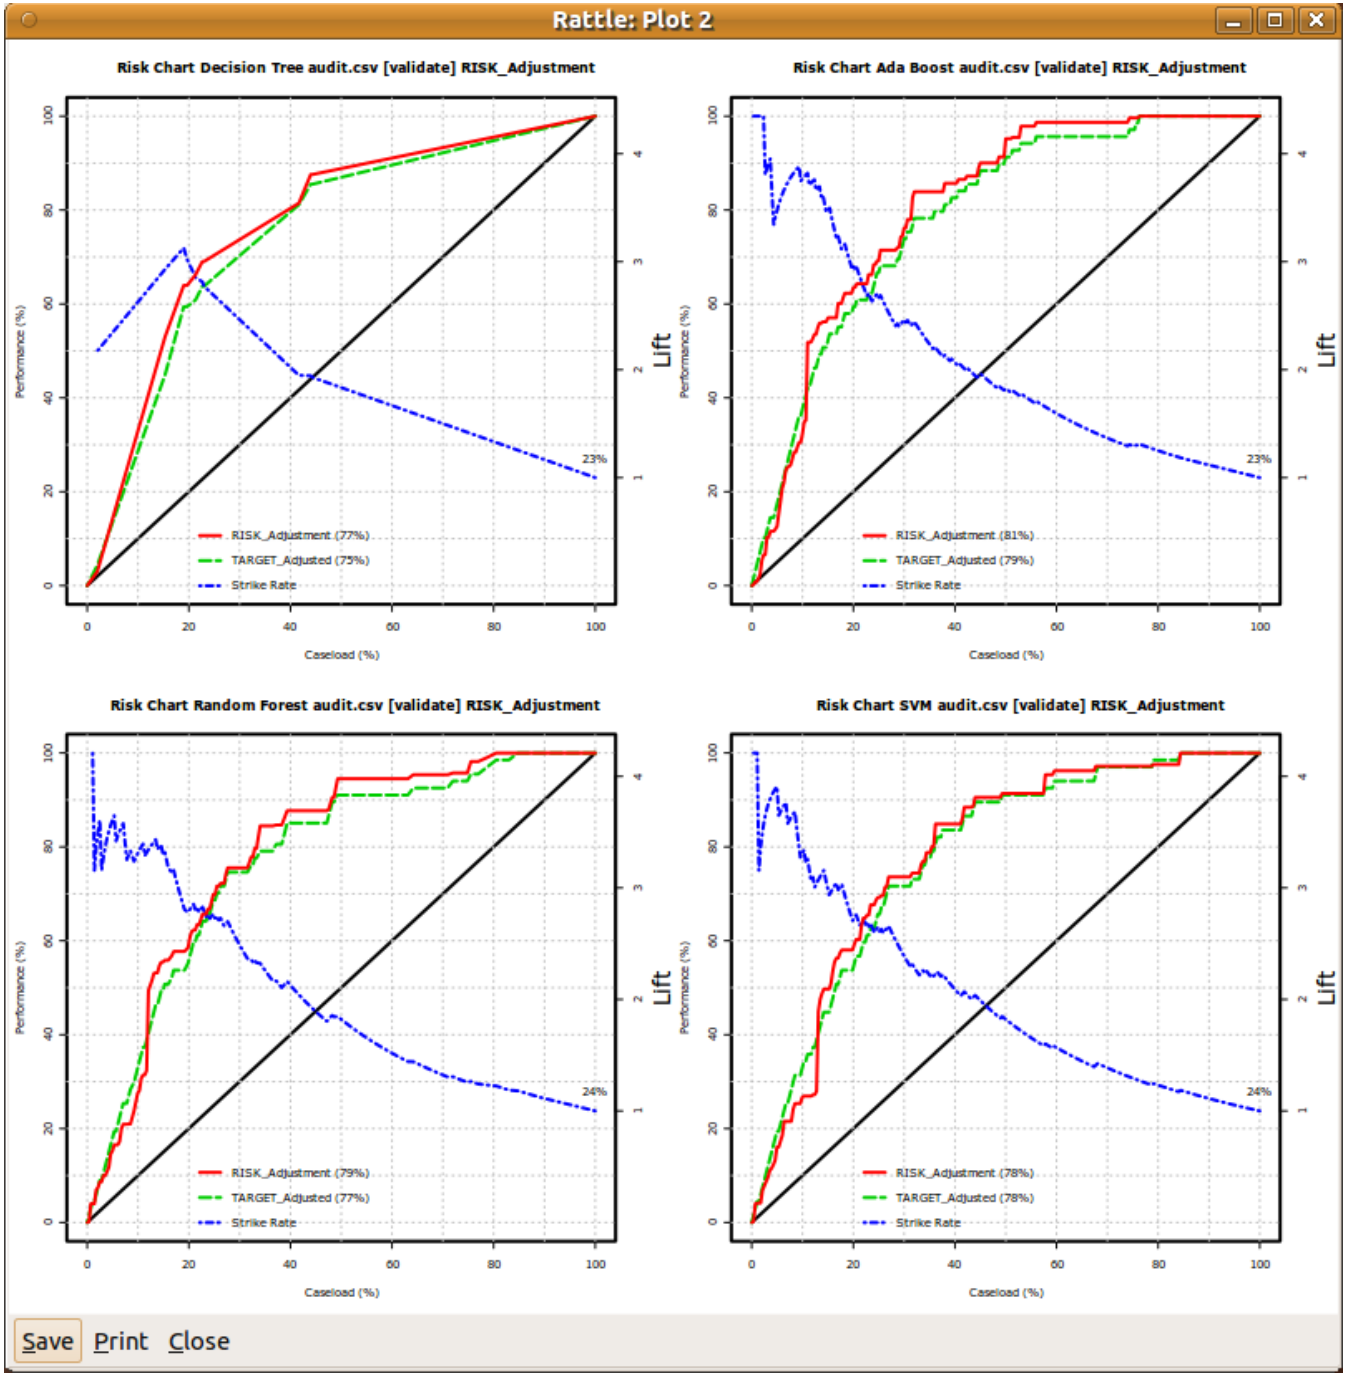

Figure 15.5: Four risk charts displayed to compare performances of multiple model builders on the audit dataset.

First Model

(example)

First model

- Once we have processed our data, we can build a model

1) Click on the Execute button
    - Rattle will notice that no dataset has been identified
2) The sample “weather” dataset will be offered
    - Click “Yes”
3) Click on the Model tab
    - This is where we tell Rattle what kind of model we want to build
4) Click on the Execute button.

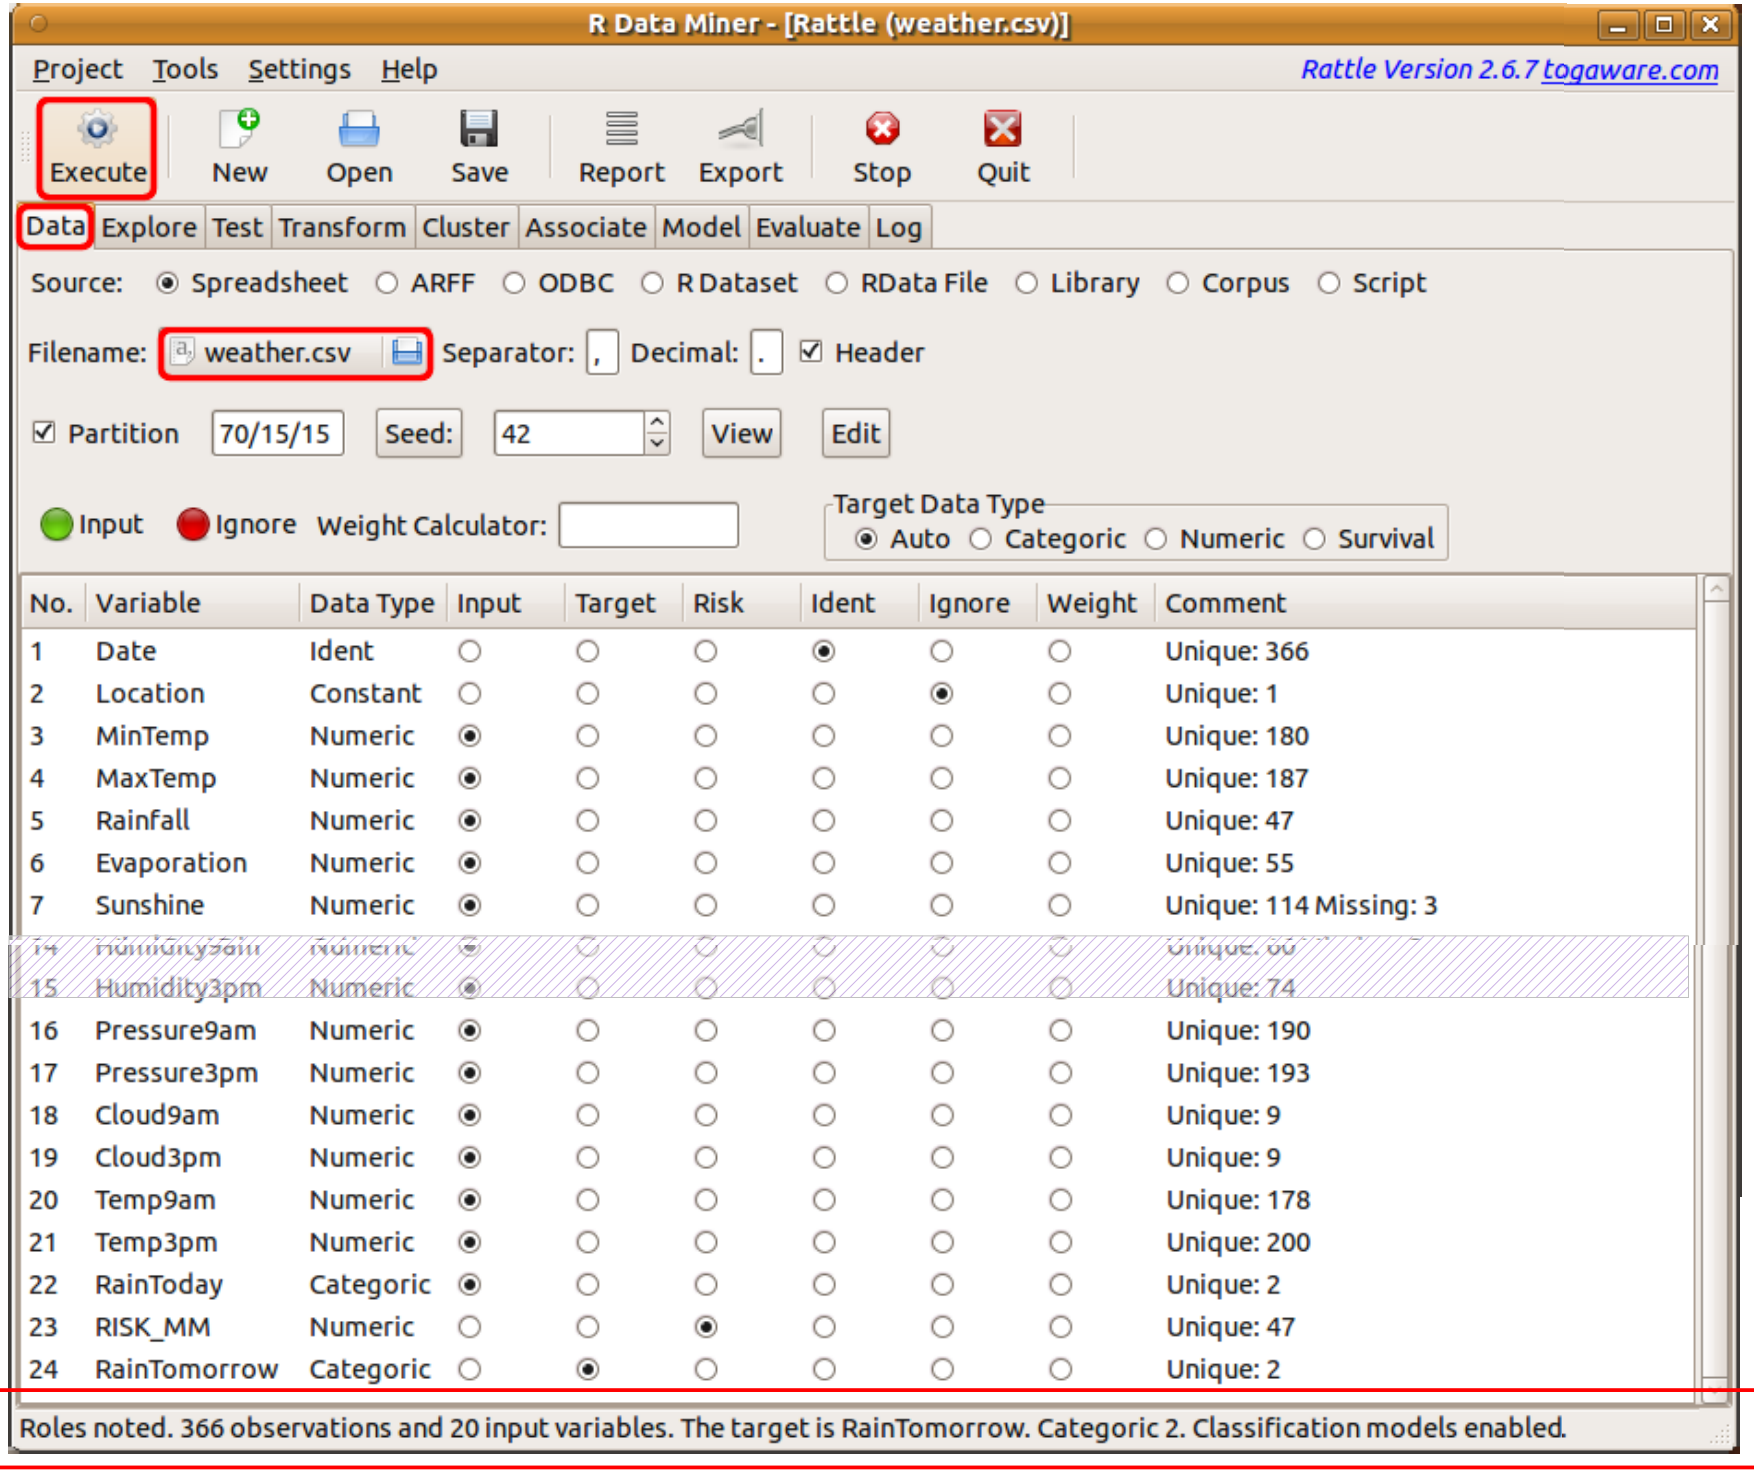
R Data Miner - [Rattle (weather.csv)]
Project Tools Settings Help
Rattle Version 2.6.7 togaware.com
Execute New Open Save Report Export Stop Quit
Data Explore Test Transform Cluster Associate Model Evaluate Log
Source: Spreadsheet ARFF ODBC R Dataset RData File Library Corpus Script
Filename: weather.csv Separator: , Decimal: . Header
Partition 70/15/15 Seed: 42 View Edit
Input Ignore Weight Calculator:
Target Data Type
Auto Categoric Numeric Survival
No. Variable Data Type Input Target Risk Ident Ignore Weight Comment
1 Date Ident Unique: 366
2 Location Constant Unique: 1
3 MinTemp Numeric Unique: 180
4 MaxTemp Numeric Unique: 187
5 Rainfall Numeric Unique: 47
6 Evaporation Numeric Unique: 55
7 Sunshine Numeric Unique: 114 Missing: 3
15 Humidity3pm Numeric Unique: 74
16 Pressure9am Numeric Unique: 190
17 Pressure3pm Numeric Unique: 193
18 Cloud9am Numeric Unique: 9
19 Cloud3pm Numeric Unique: 9
20 Temp9am Numeric Unique: 178
21 Temp3pm Numeric Unique: 200
22 RainToday Categoric Unique: 2
23 RISK_MM Numeric Unique: 47
24 RainTomorrow Categoric Unique: 2
Roles noted. 366 observations and 20 input variables. The target is RainTomorrow. Categoric 2. Classification models enabled.

Building a Model

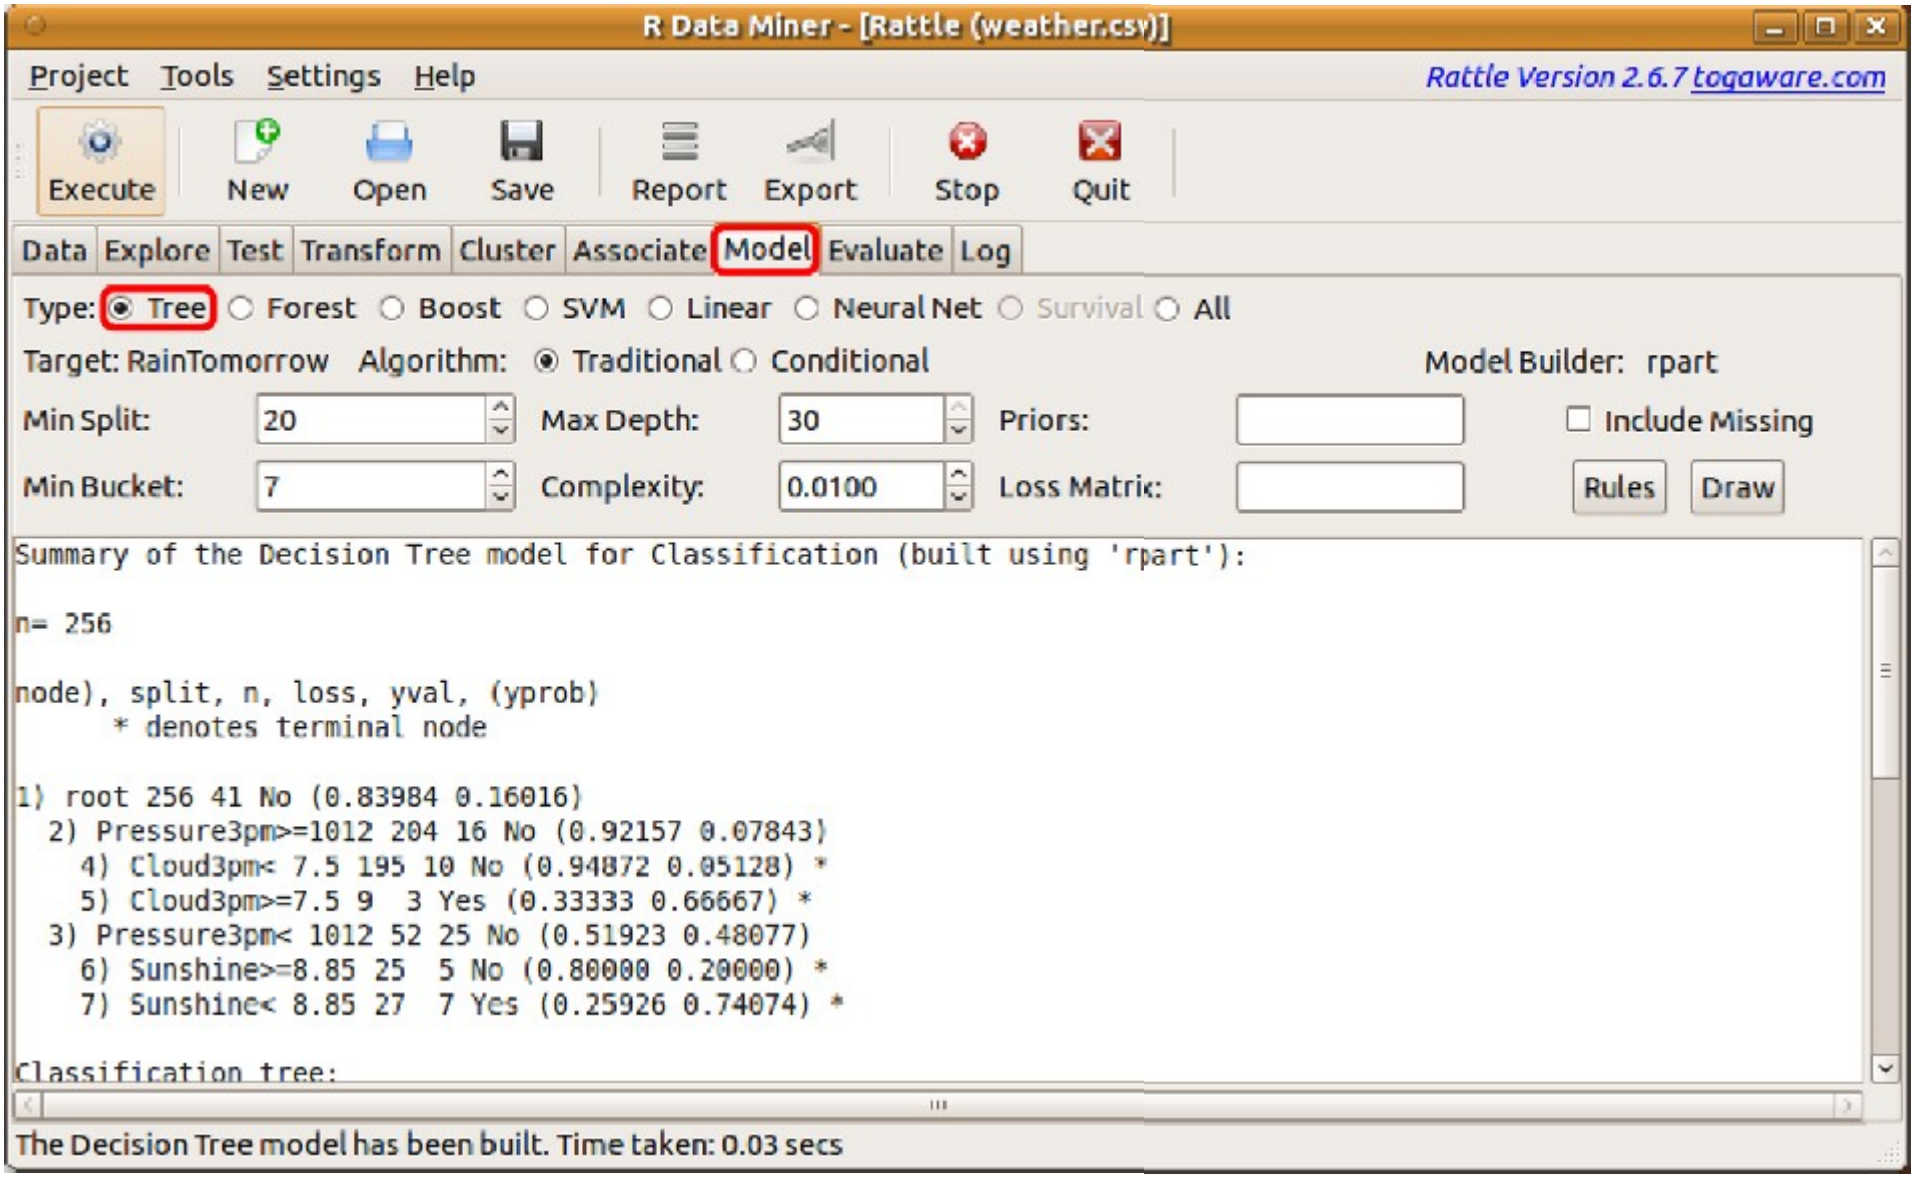

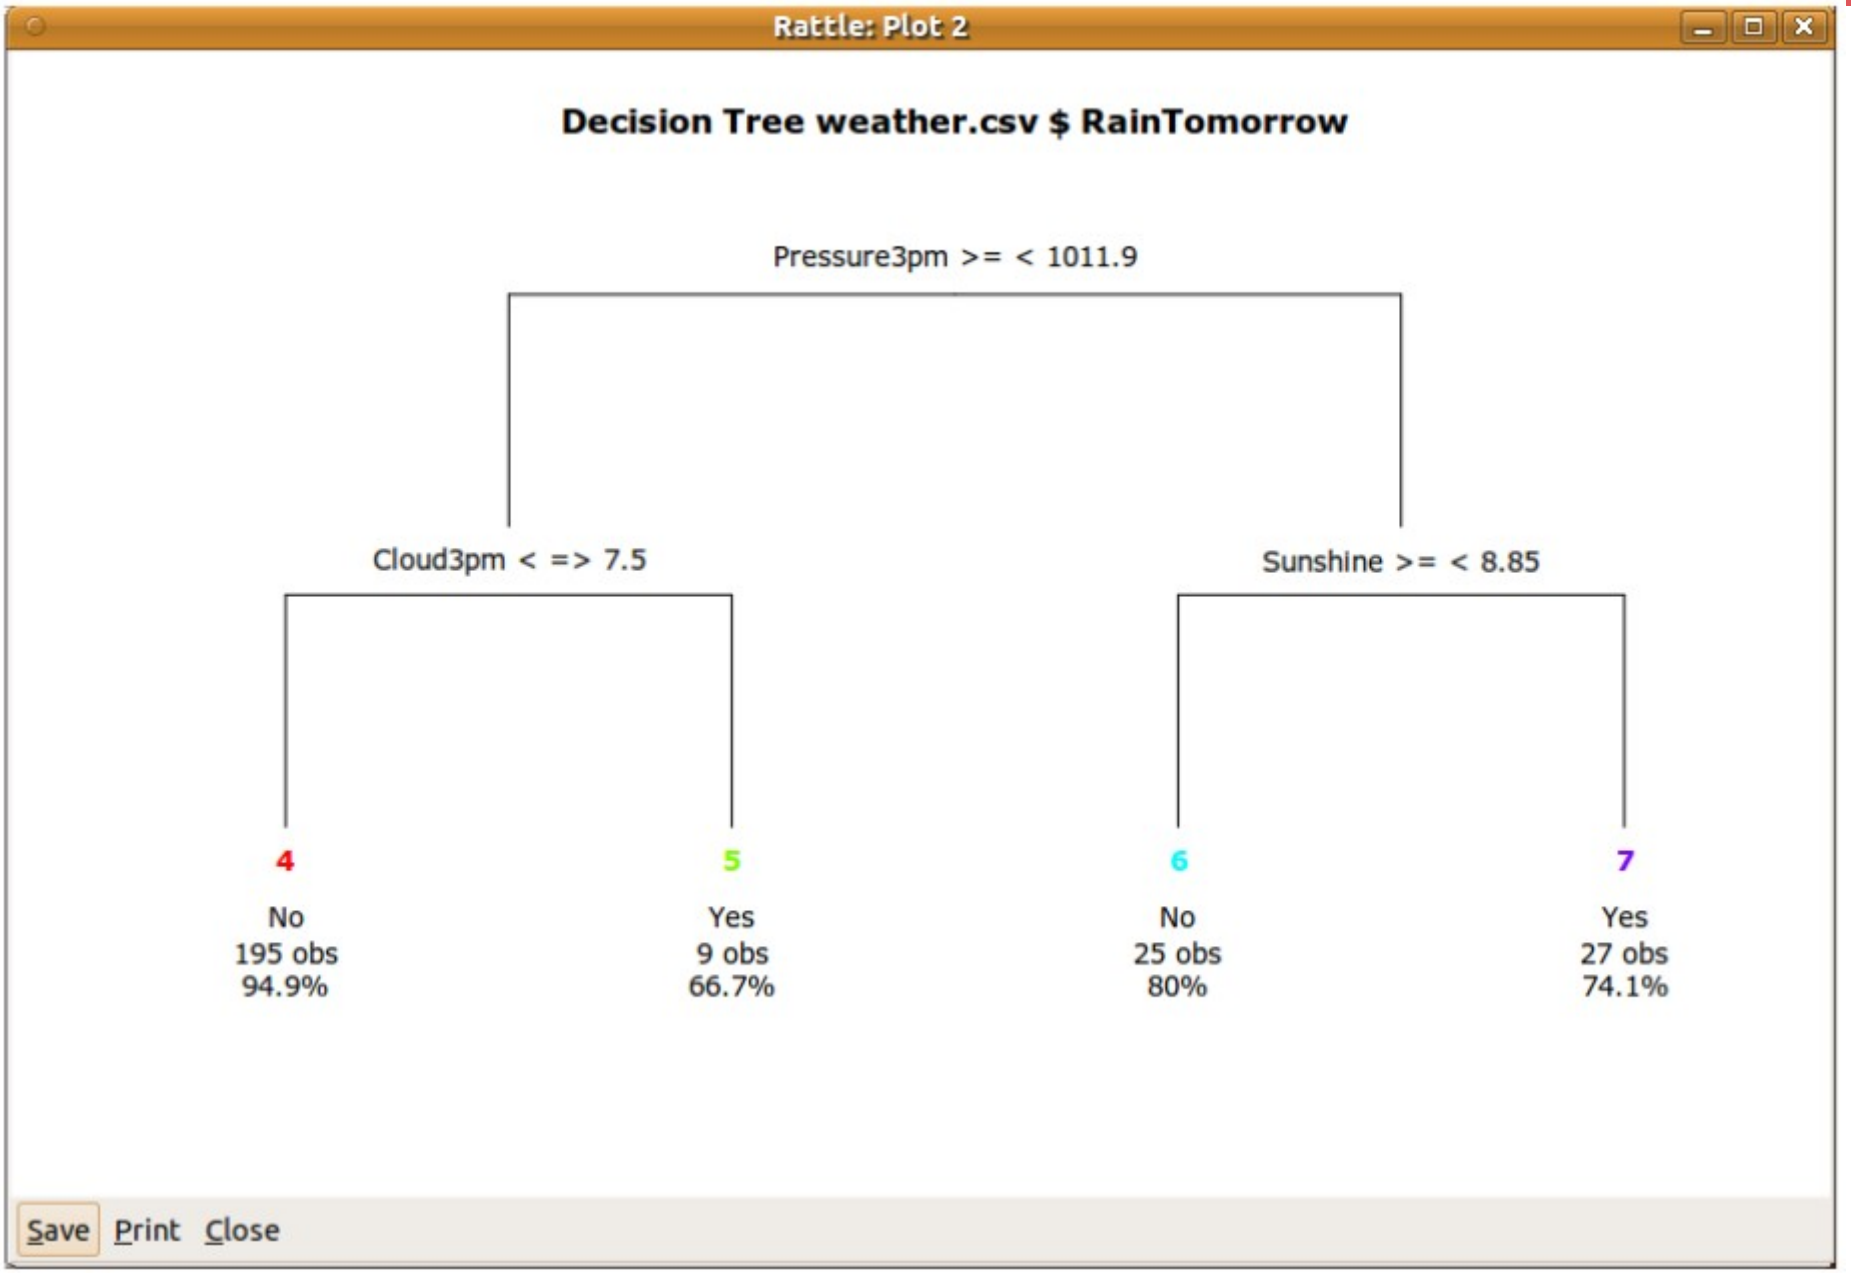
Rattle: Plot 2
Decision Tree weather.csv $ RainTomorrow
Pressure3pm >= < 1011.9
Cloud3pm < => 7.5
Sunshine >= < 8.85
4
No
195 obs
94.9%
5
Yes
9 obs
66.7%
6
No
25 obs
80%
7
Yes
27 obs
74.1%
Save
Print
Close

Figure 2.5

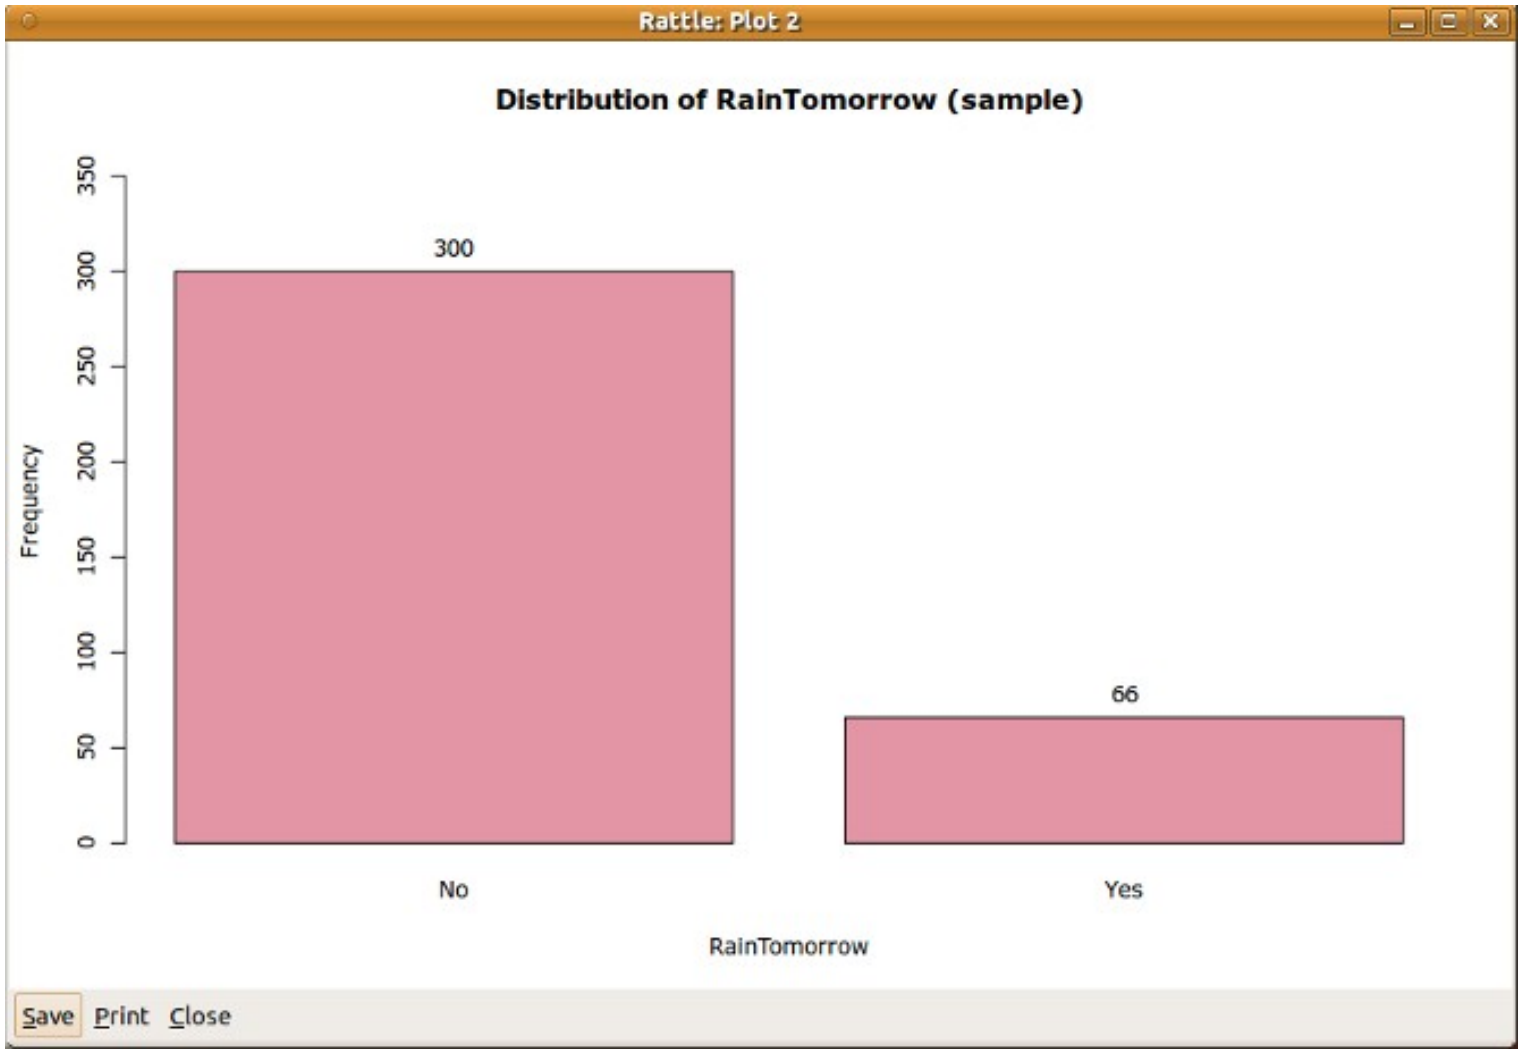

Figure 2.6: The target variable, RainTomorrow, is skewed, with Yes being quite under-represented

Appendix

Installation

Install R

```
www.r-project.org
```

```
Start R
```

Install Rattle

```
> install.packages("rattle")
```

Load rattle into the R library

```
> library(rattle)
> rattle()
```

Articles on Rattle

- **Rattle: A Data Mining GUI for R**
- http://journal.r-project.org/archive/2009-2/RJournal_2009-2_Williams.pdf

- **Getting started with data mining in R using Rattle**
- http://techpad.co.uk/content.php?sid=240

R Resources (sample)

- **Data Manipulation with R** (Spector, 2008) – covers basic data structures, reading and writing data, subscripting, manipulating, aggregating, and reshaping data
- **Introductory Statistics with R** (Dalgaard, 2008) - good introduction to statistics using R.
- **Modern Applied Statistics with S** (Venables and Ripley, 2002 - an extensive introduction to statistics using R.
- **Data Analysis and Graphics Using R** (Maindonald and Braun, 2007) - excellent practical coverage of many aspects of exploring and modelling data using R
- **The Elements of Statistical Learning** (Hastie et al., 2009) is a more mathematical treatise, covering all of the machine learning techniques discussed in this book in quite some mathematical depth.
- **R for SAS and SPSS Users** (Muenchen, 2008) is an excellent choice
- **Lattice: Multivariate Data Visualization with R** (Sarkar, 2008) - covers
- the extensive capabilities of one of the graphics/plotting packages available for R.
- **ggplot2: Elegant Graphics for Data Analysis** (Wickham, 2009) - newer graphics framework is detailed
- **Bivand et al.** (2008) cover applied spatial data analysis,
- **Kleiber and Zeileis** (2008) cover applied econometrics
- **Cowpertwait and Metcalfe** (2009) cover time series

Agenda

- Intro
- Data mining
- Rattle installation
- Rattle workflow
- Appendix

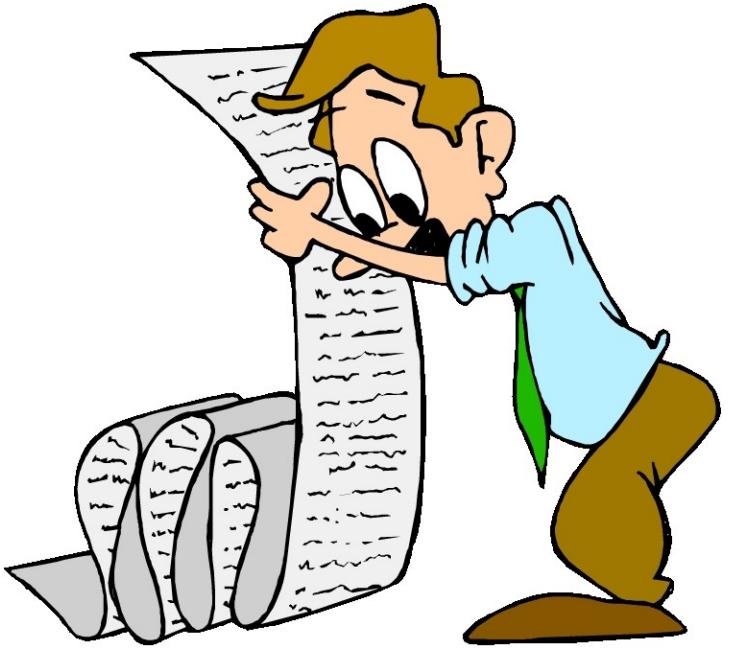

“Data Mining with Rattle and R”

DIT Analytics Club
07 March, 2013
“Peadar Kearney's”, Dame Street

Colman McMahon
colman@thedata.co

LinkedIn
www.linkedin.com/in/colmanmcmahon